

Classic Solutions Architecture

Section Introduction

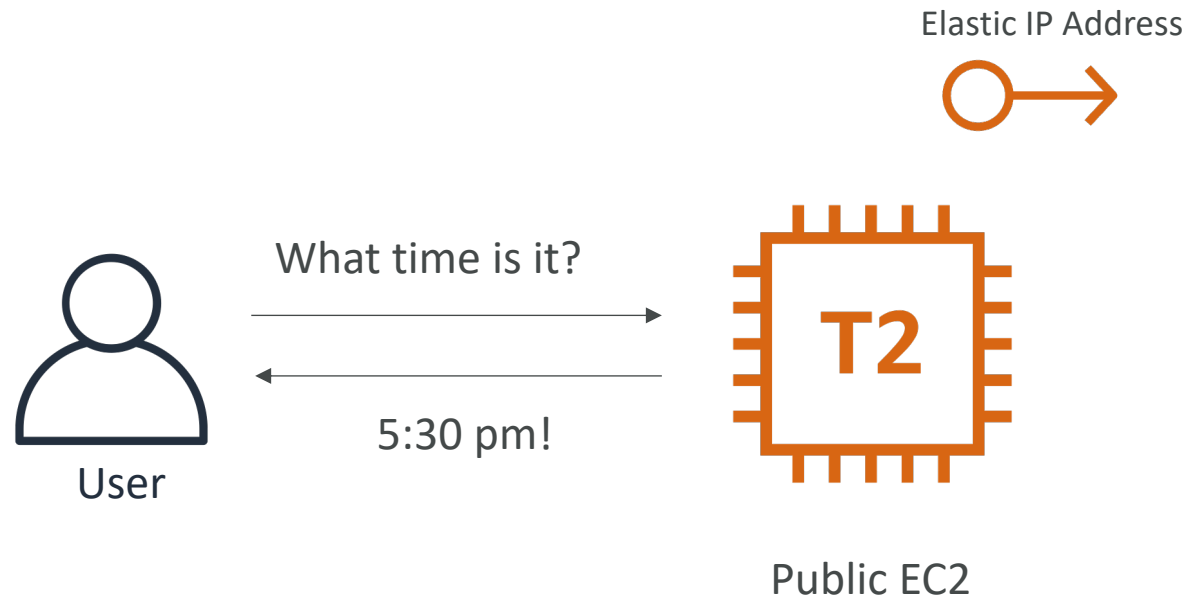
- These solutions architectures are the best part of this course
- Let's understand how all the technologies we've seen work together
- This is a section you need to be 100% comfortable with
- We'll see the progression of a Solution's architect mindset through many sample case studies:
 - WhatIsTheTime.Com
 - MyClothes.Com
 - MyWordPress.Com
 - Instantiating applications quickly
 - Beanstalk

Stateless Web App: WhatIsTheTime.com

- WhatIsTheTime.com allows people to know what time it is
- We don't need a database
- We want to start small and can accept downtime
- We want to fully scale vertically and horizontally, no downtime
- Let's go through the Solutions Architect journey for this app
- Let's see how we can proceed!

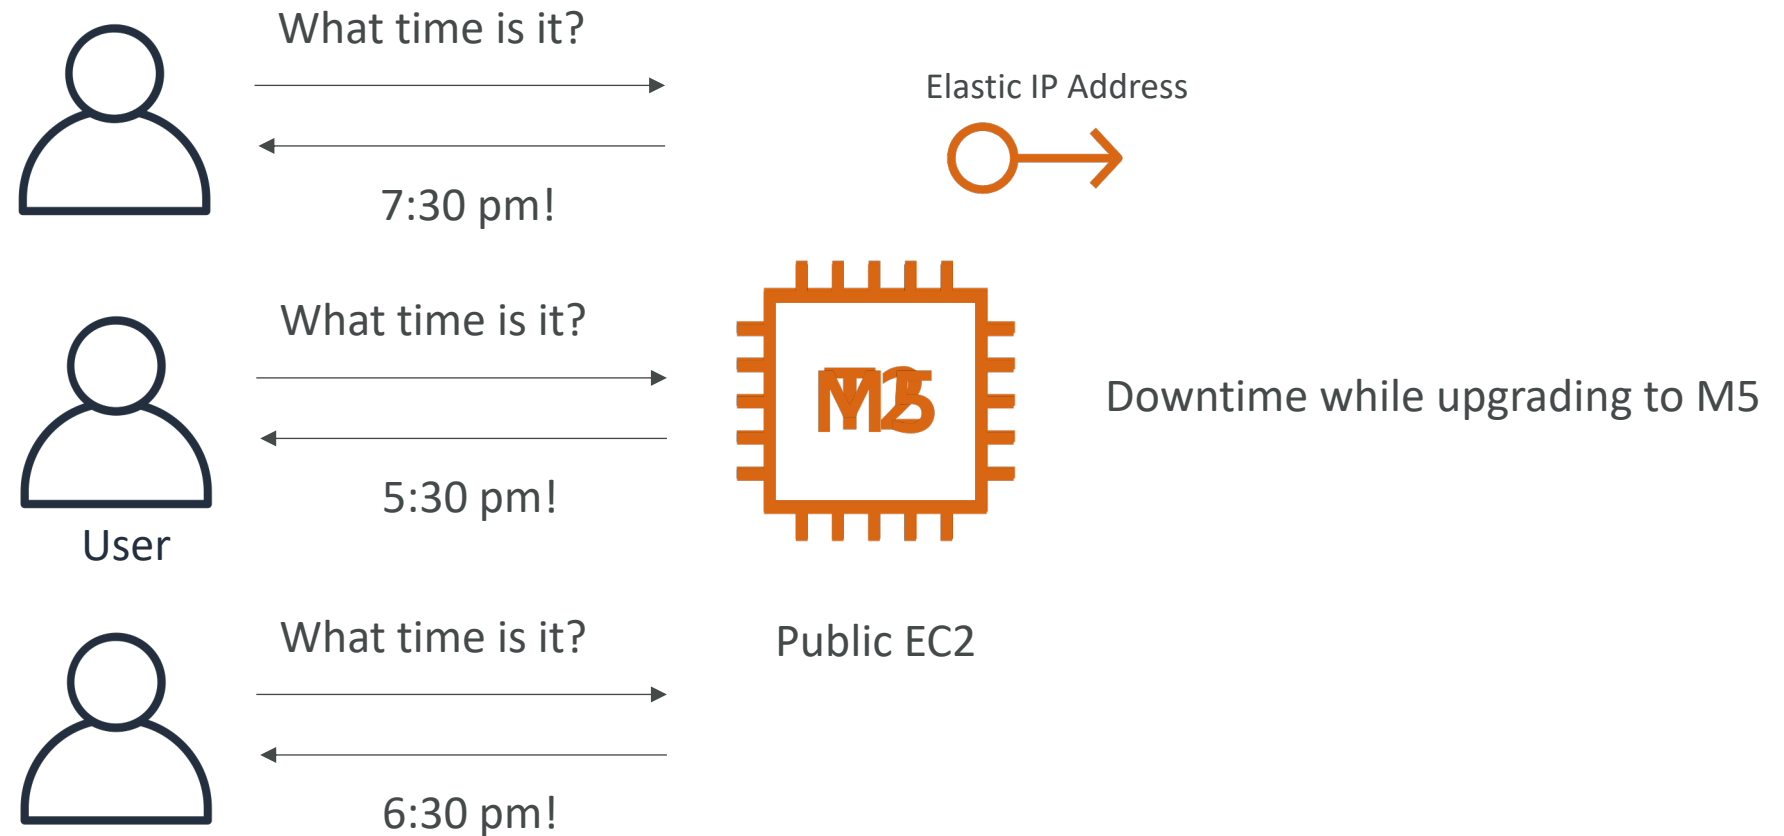
Stateless web app: What time is it?

Starting simple



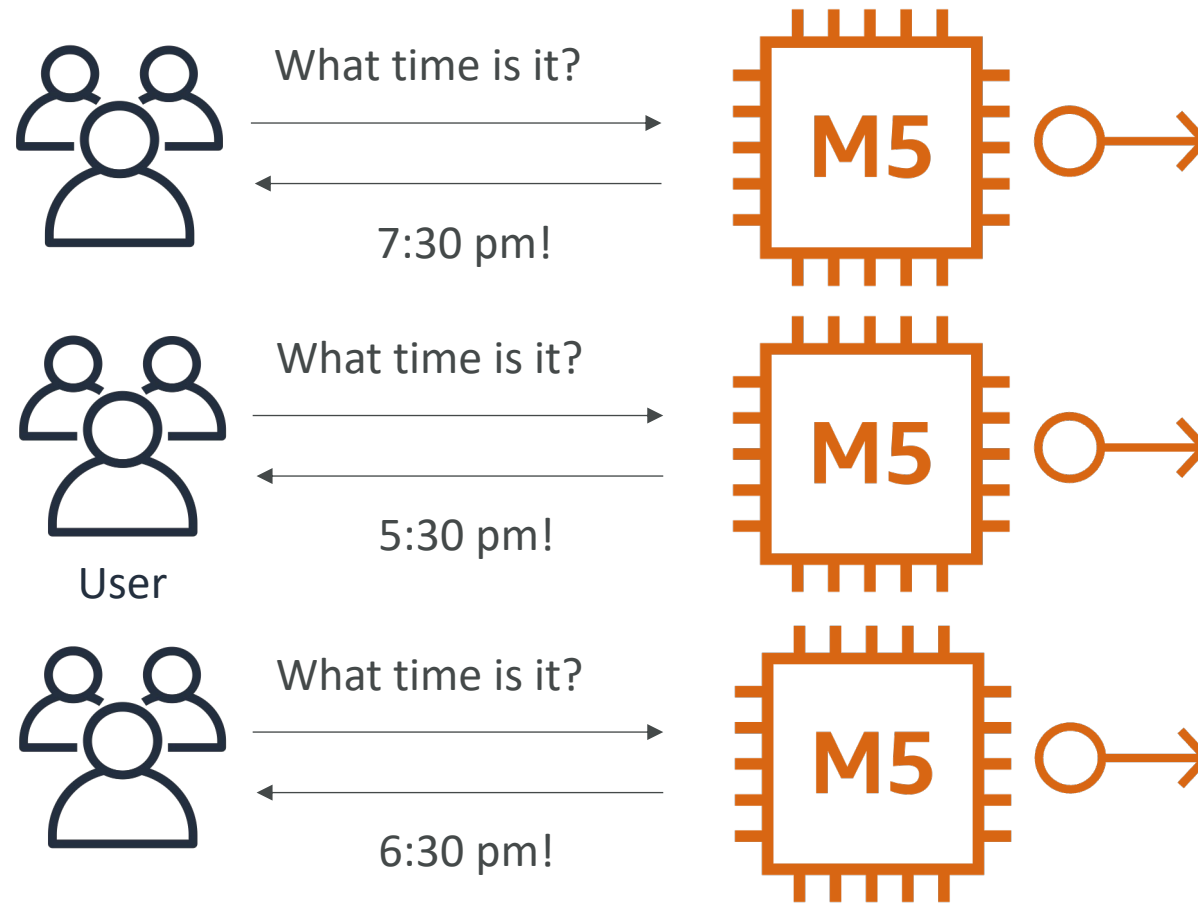
Stateless web app: What time is it?

Scaling vertically



Stateless web app: What time is it?

Scaling horizontally



Stateless web app: What time is it?

Scaling horizontally

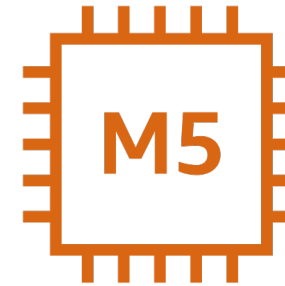
DNS Query
For api.whatisthetime.com
A Record
TTL 1 hour



What time is it?



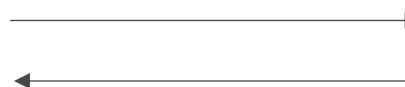
7:30 pm!



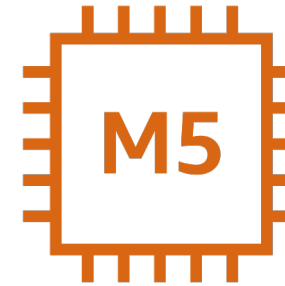
Public EC2 instance,
No Elastic IP



What time is it?



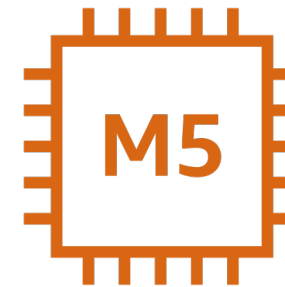
5:30 pm!



What time is it?



6:30 pm!



Stateless web app: What time is it?

Scaling horizontally, adding and removing instances

DNS Query
For api.whatisthetime.com
A Record
TTL 1 hour



What time is it?

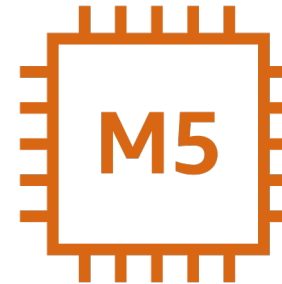
7:30 pm!

INSTANCE IS GONE!



What time is it?

5:30 pm!

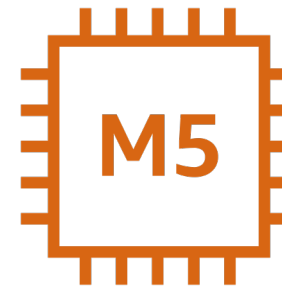


Public EC2 instance,
No Elastic IP



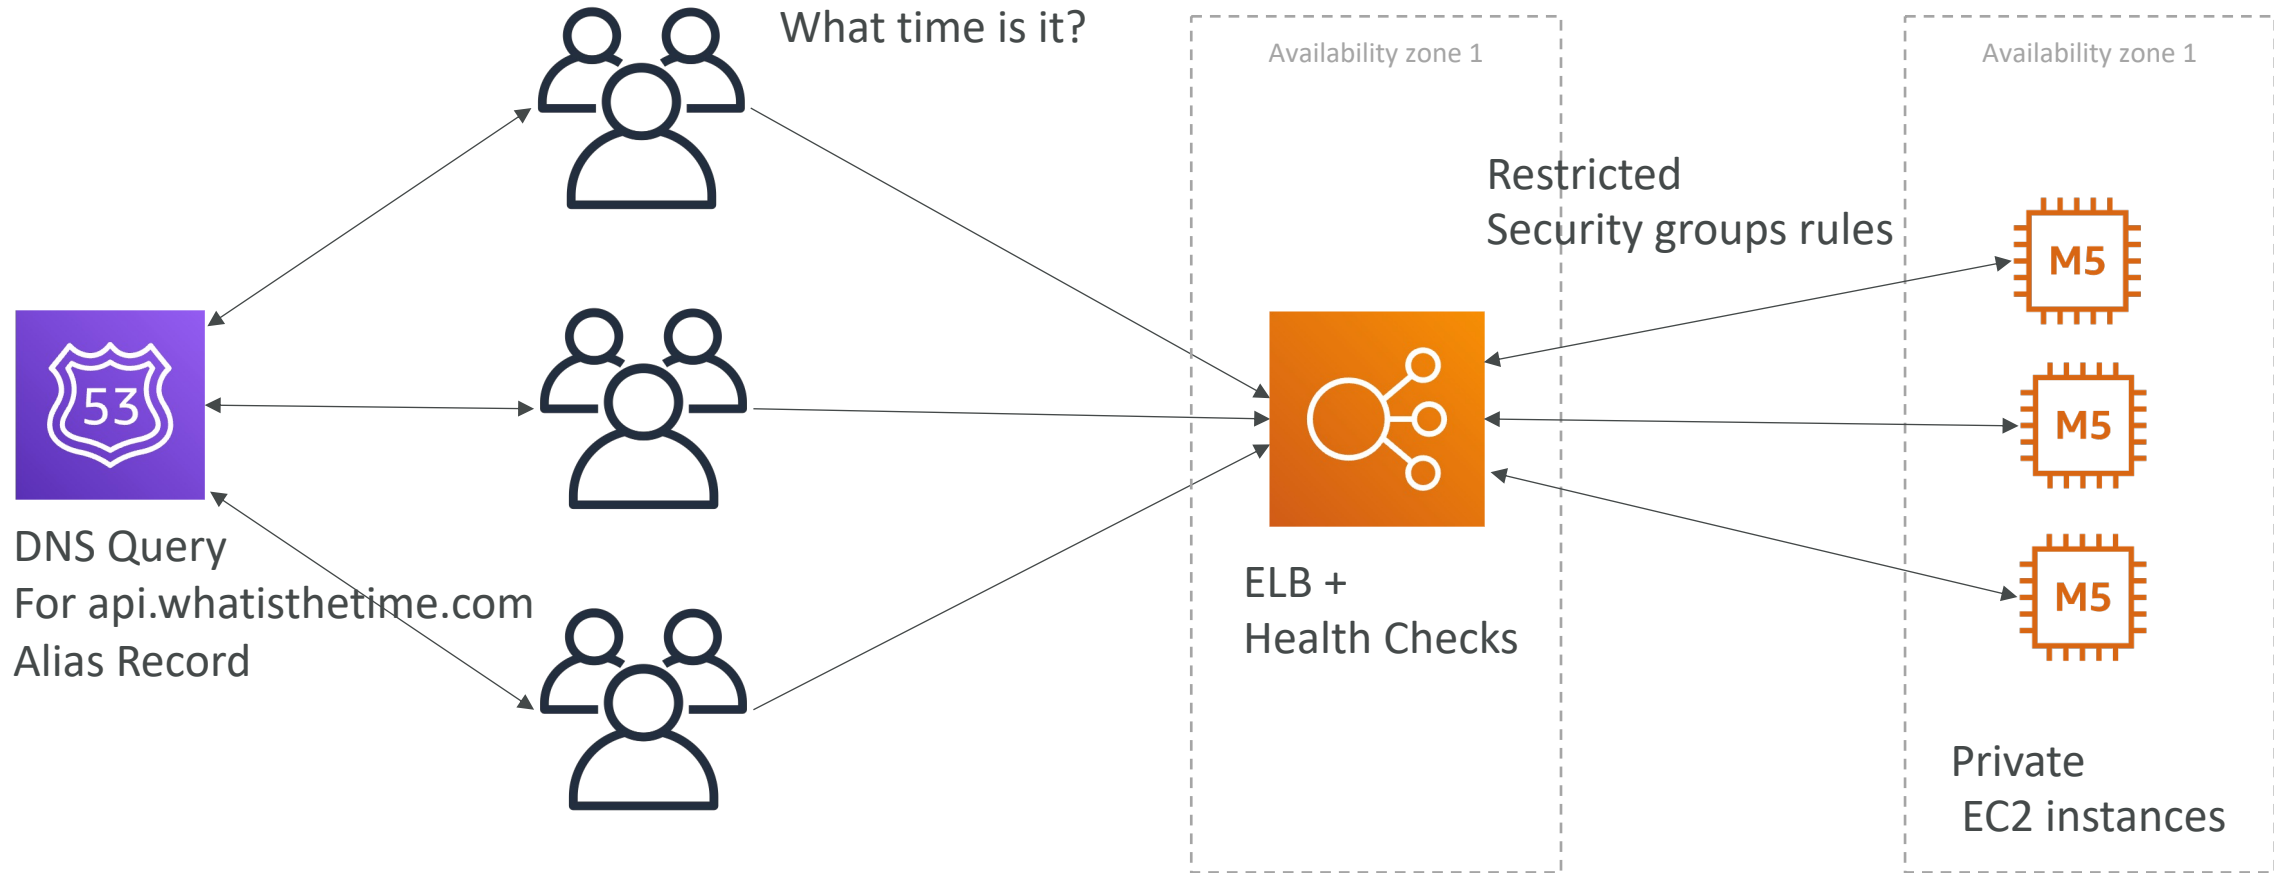
What time is it?

6:30 pm!



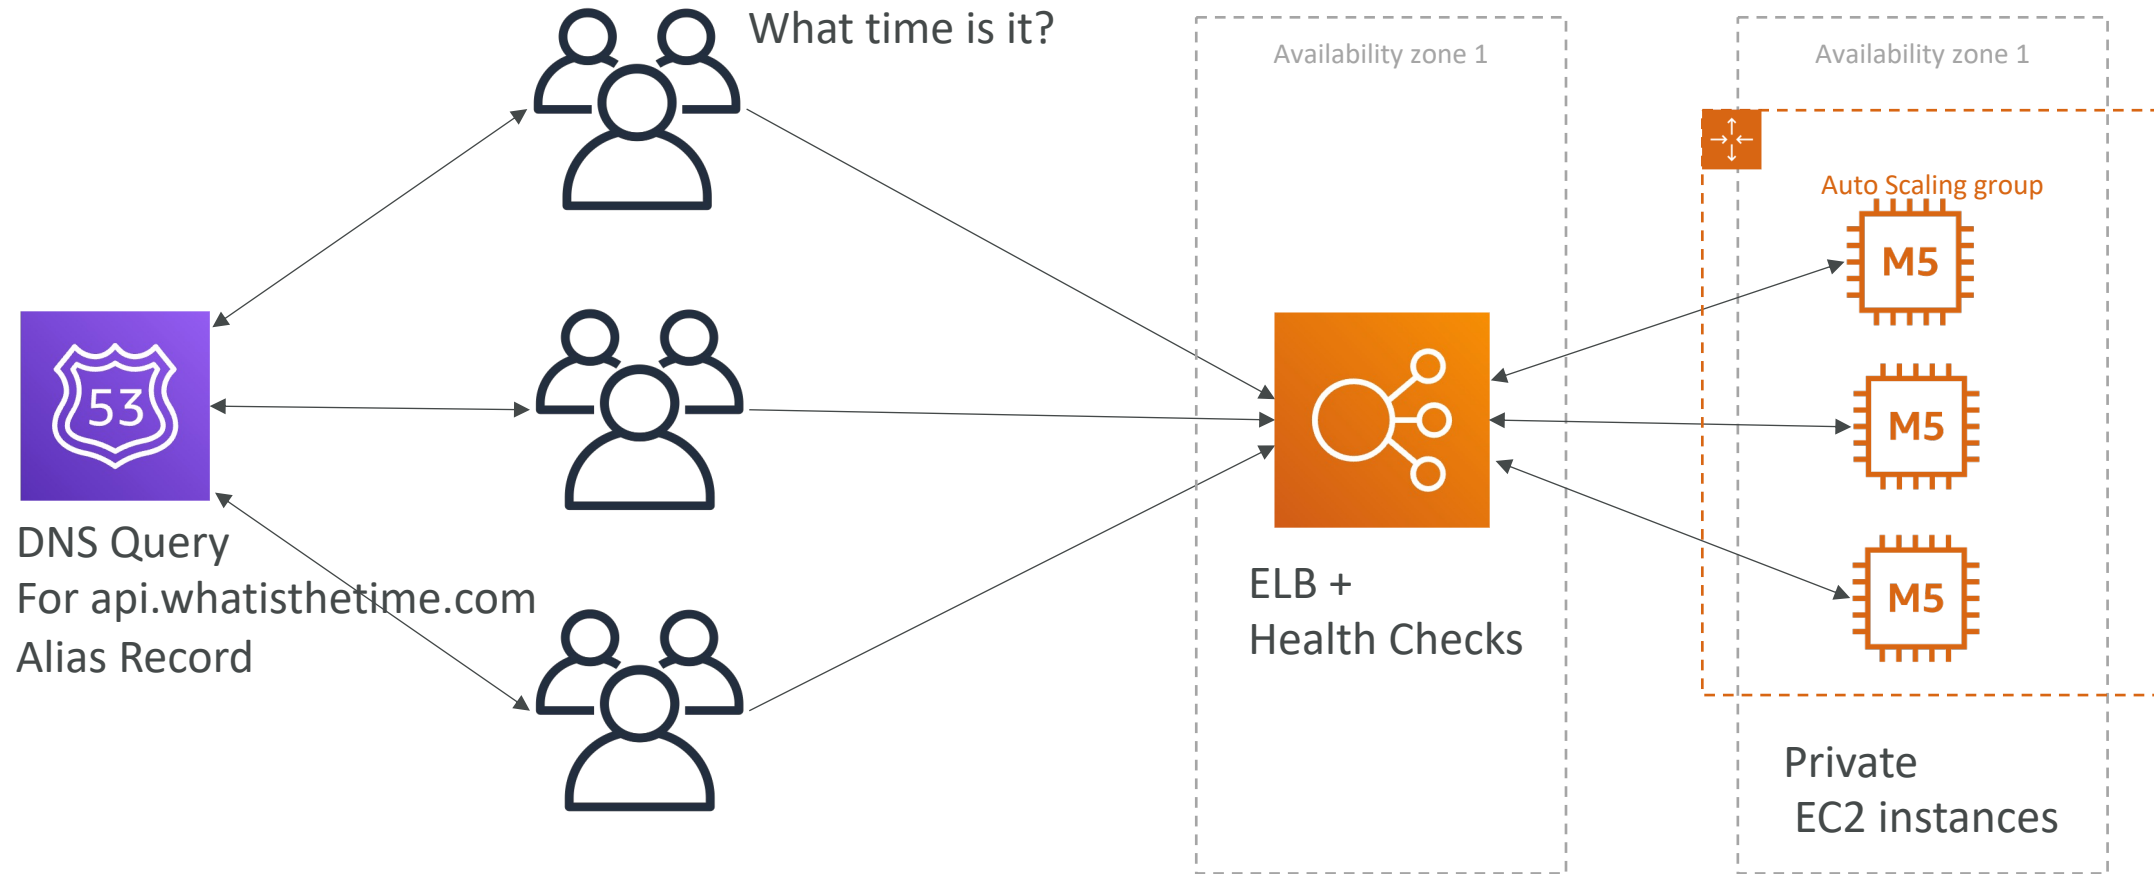
Stateless web app: What time is it?

Scaling horizontally, with a load balancer



Stateless web app: What time is it?

Scaling horizontally, with an auto-scaling group

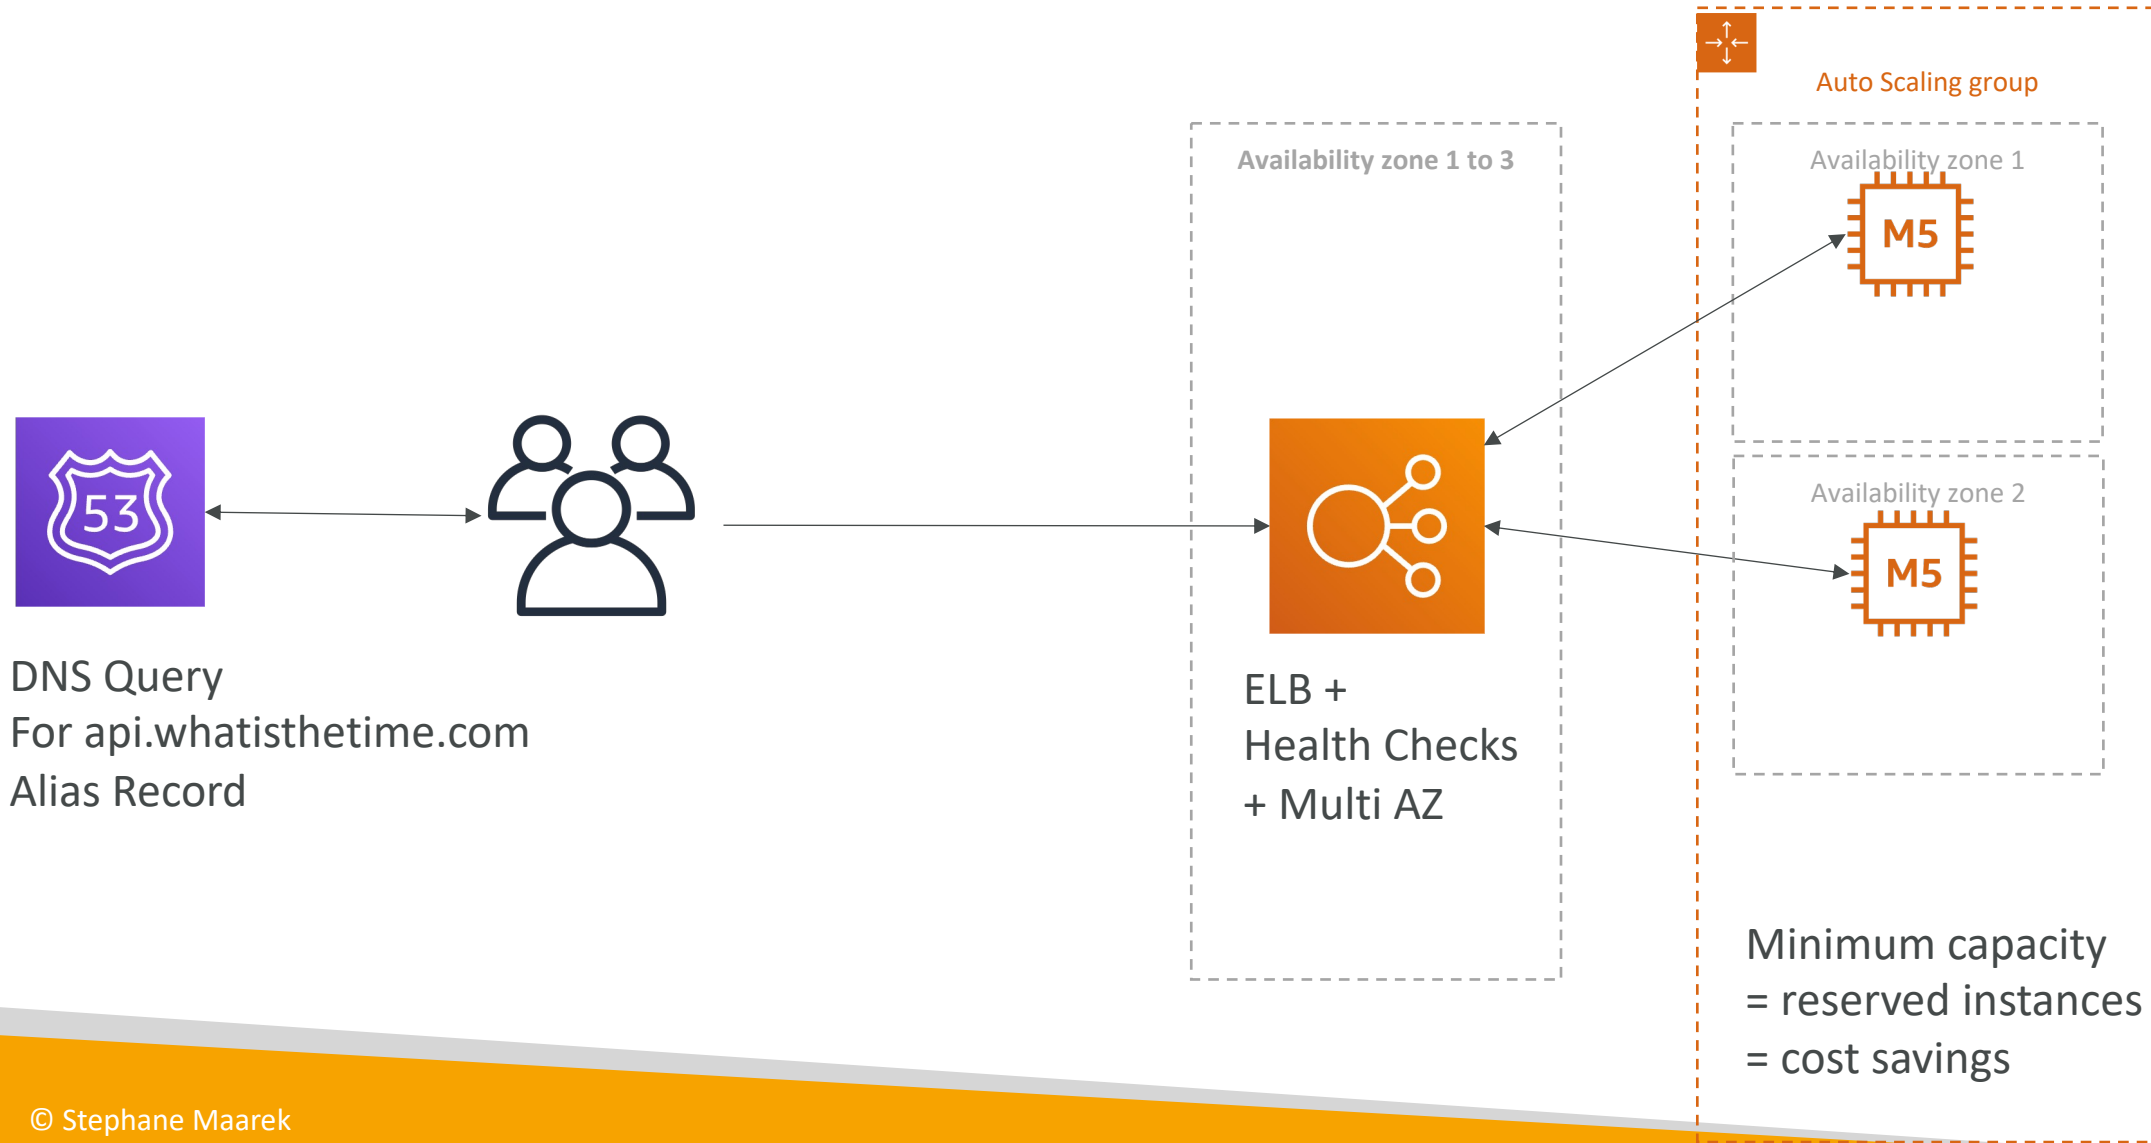


Stateless web app: What time is it?

Making our app multi-AZ



Minimum 2 AZ => Let's reserve capacity



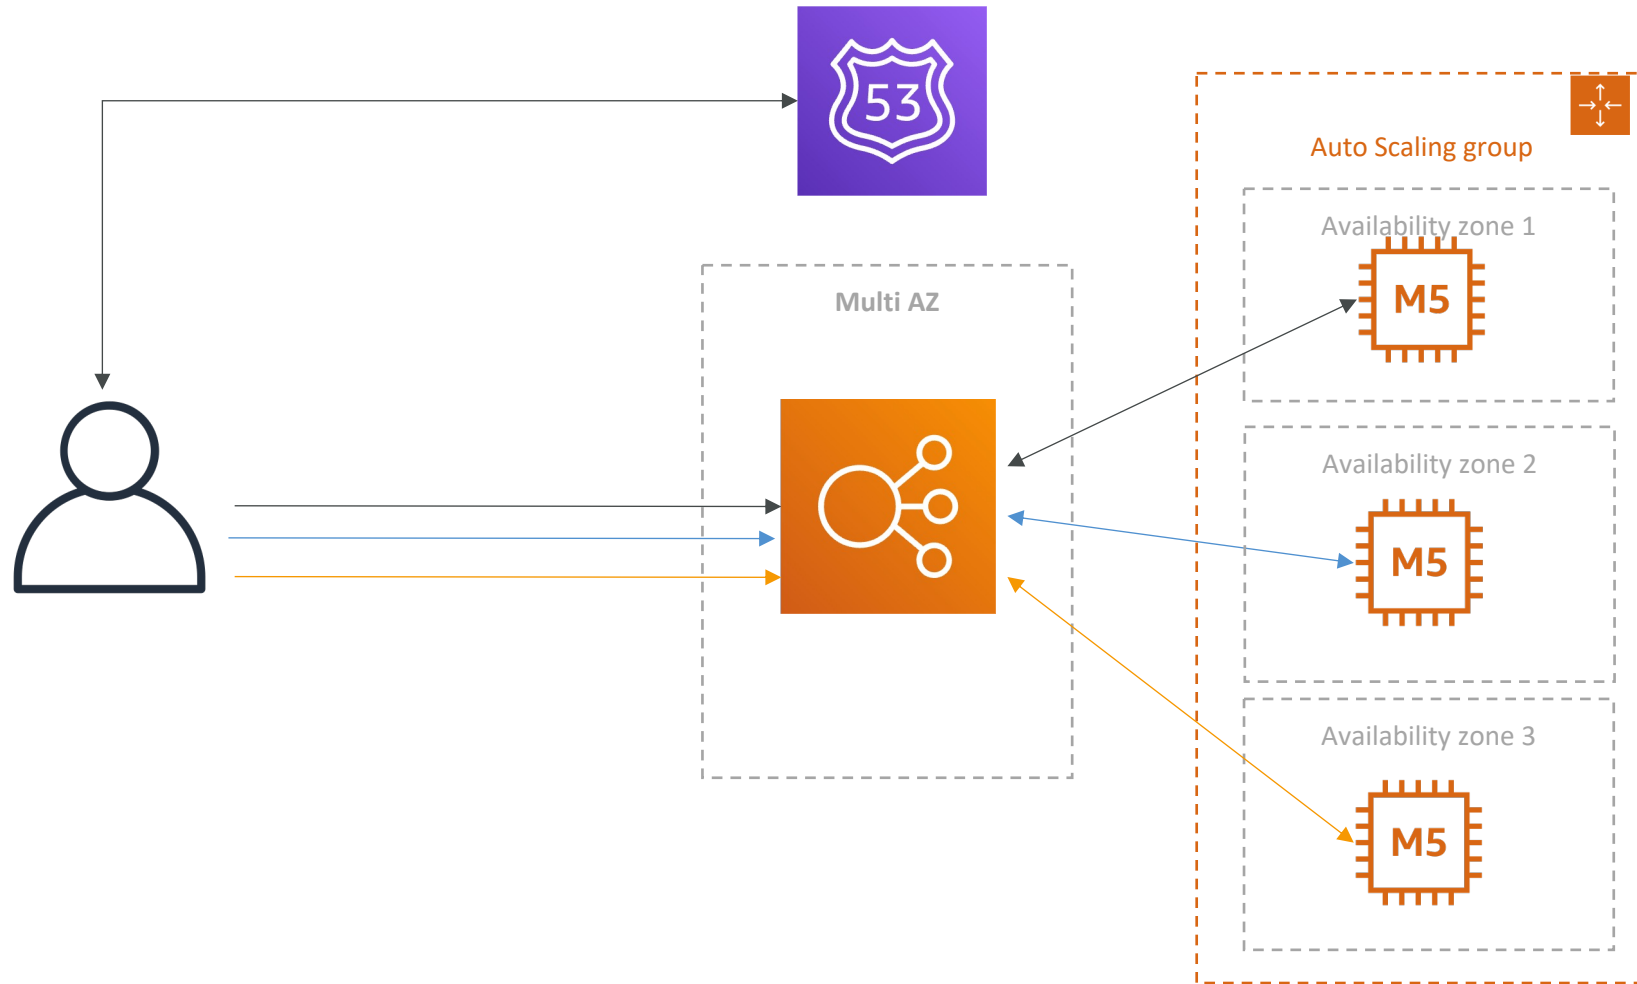
In this lecture we've discussed...

- Public vs Private IP and EC2 instances
- Elastic IP vs Route 53 vs Load Balancers
- Route 53 TTL, A records and Alias Records
- Maintaining EC2 instances manually vs Auto Scaling Groups
- Multi AZ to survive disasters
- ELB Health Checks
- Security Group Rules
- Reservation of capacity for costing savings when possible

Stateful Web App: MyClothes.com

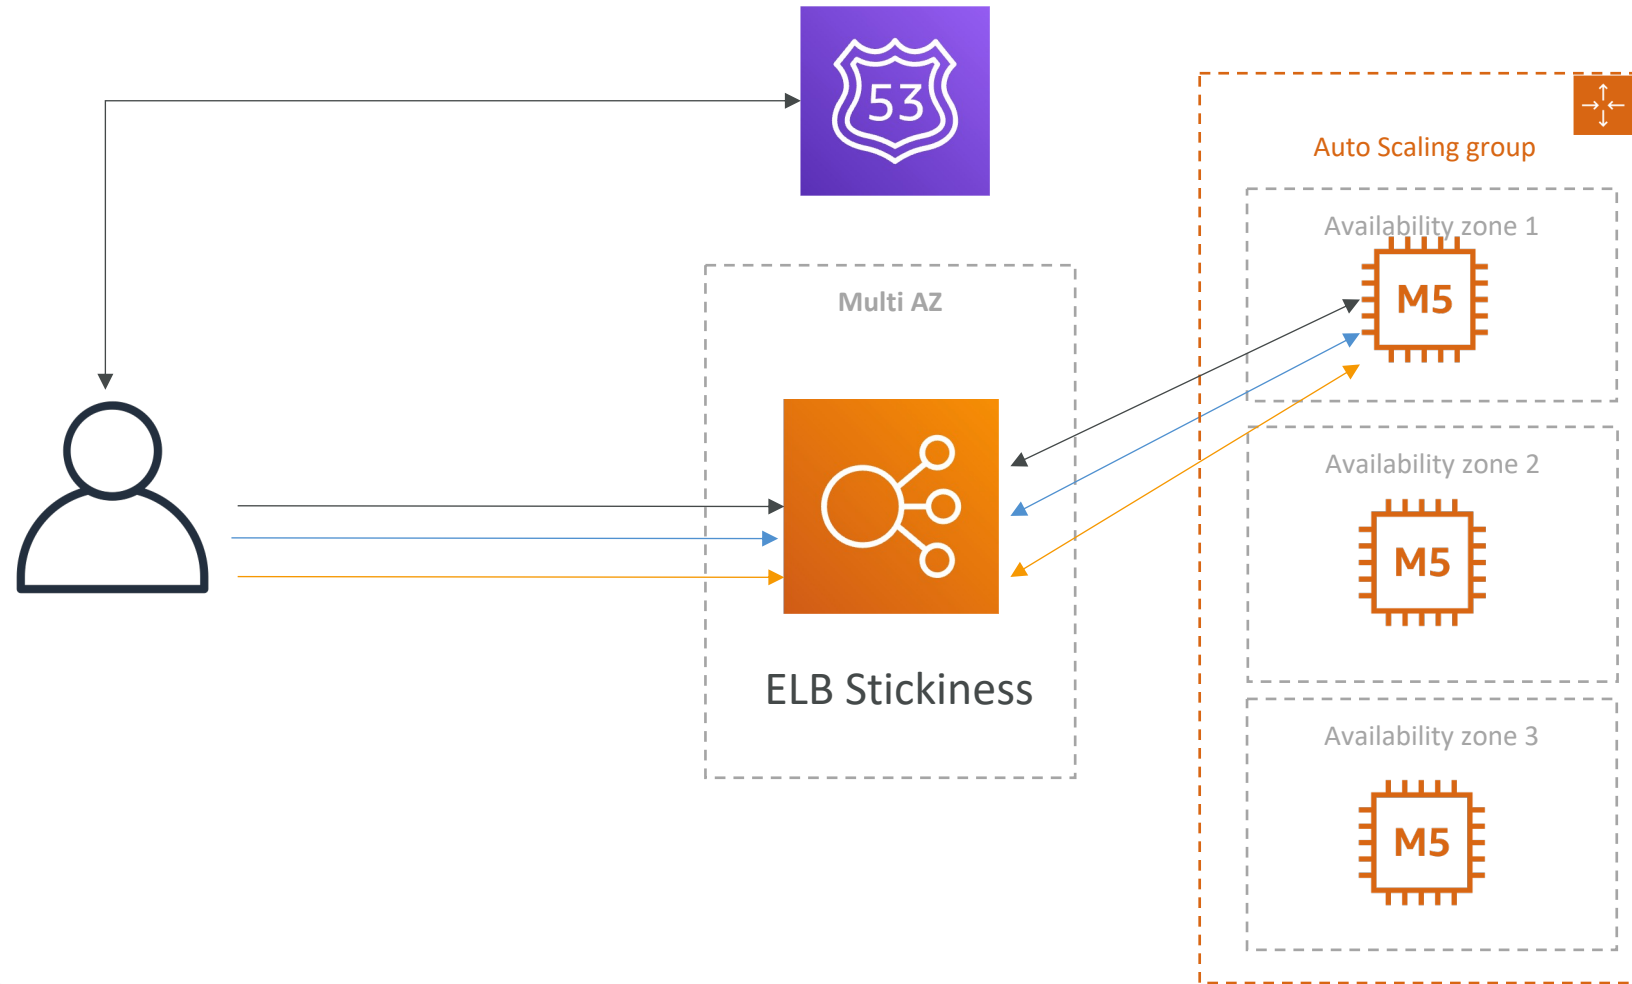
- MyClothes.com allows people to buy clothes online.
- There's a shopping cart
- Our website is having hundreds of users at the same time
- We need to scale, maintain horizontal scalability and keep our web application as stateless as possible
- Users should not lose their shopping cart
- Users should have their details (address, etc) in a database
- Let's see how we can proceed!

Stateful Web App: MyClothes.com



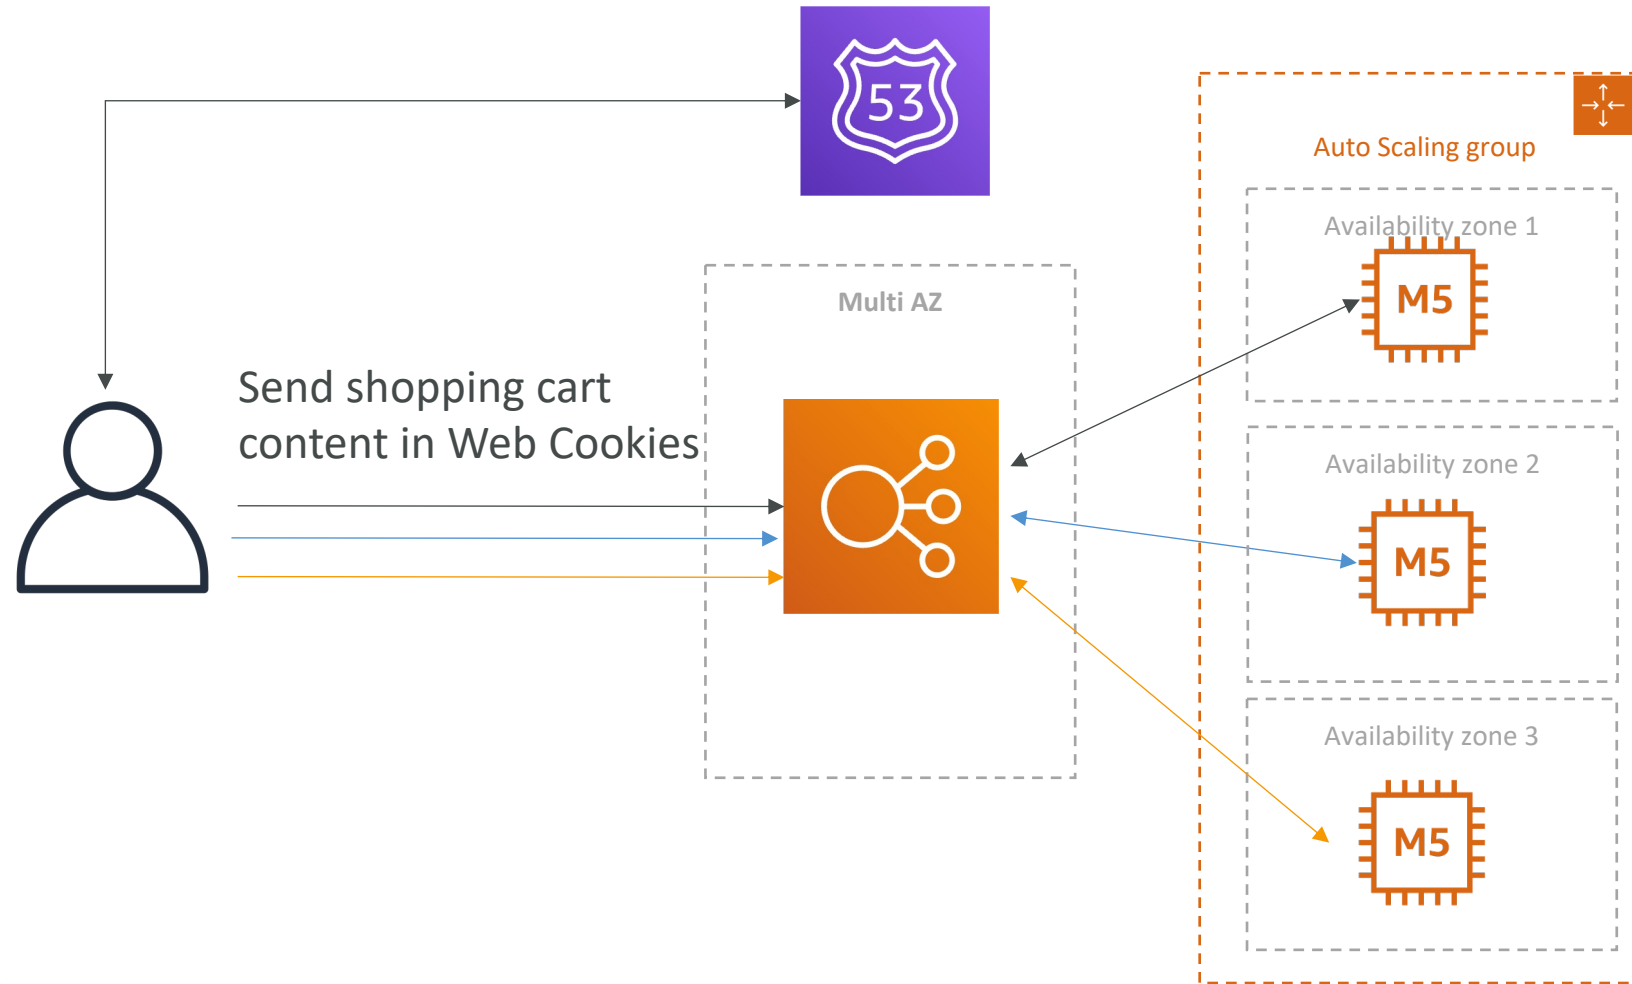
Stateful Web App: MyClothes.com

Introduce Stickiness (Session Affinity)



Stateful Web App: MyClothes.com

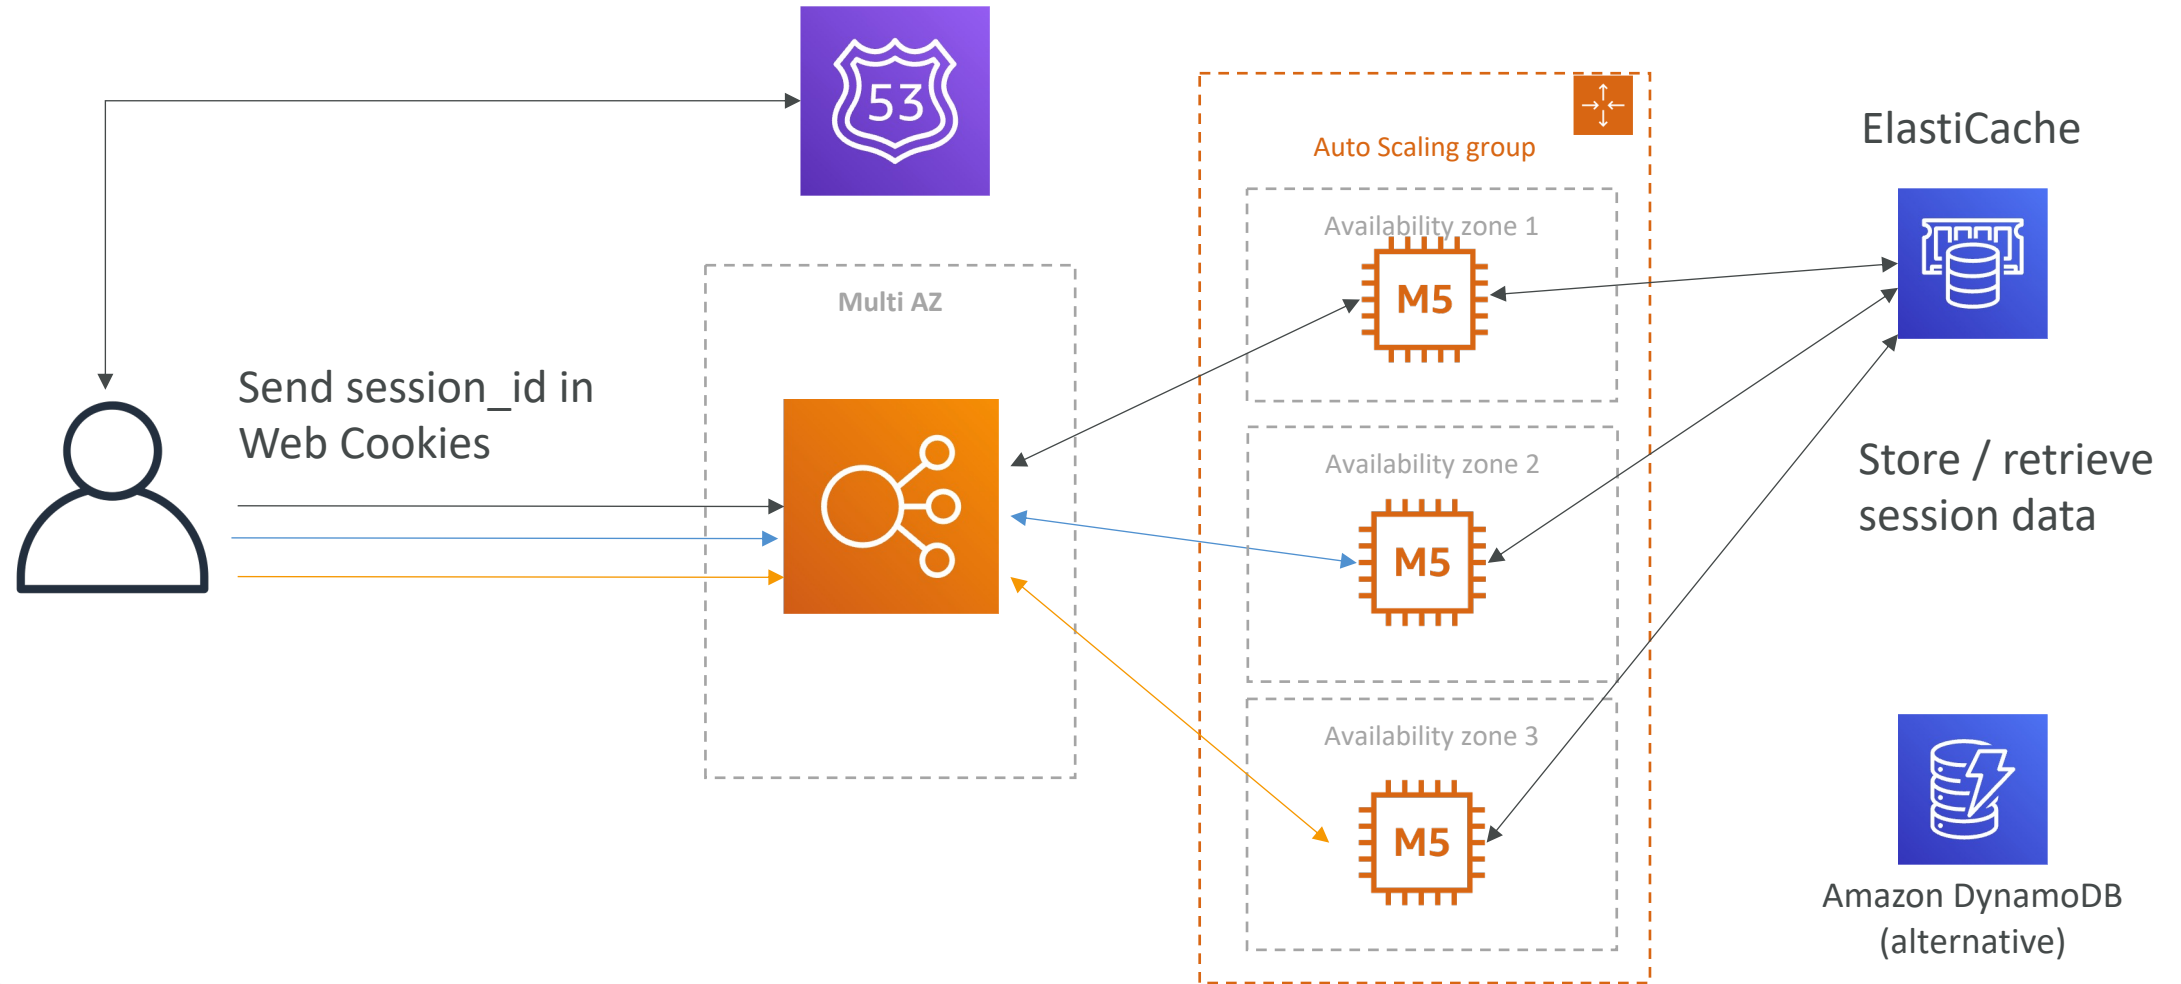
Introduce User Cookies



Stateless
HTTP requests are heavier
Security risk
(cookies can be altered)
Cookies must be validated
Cookies must be less than 4KB

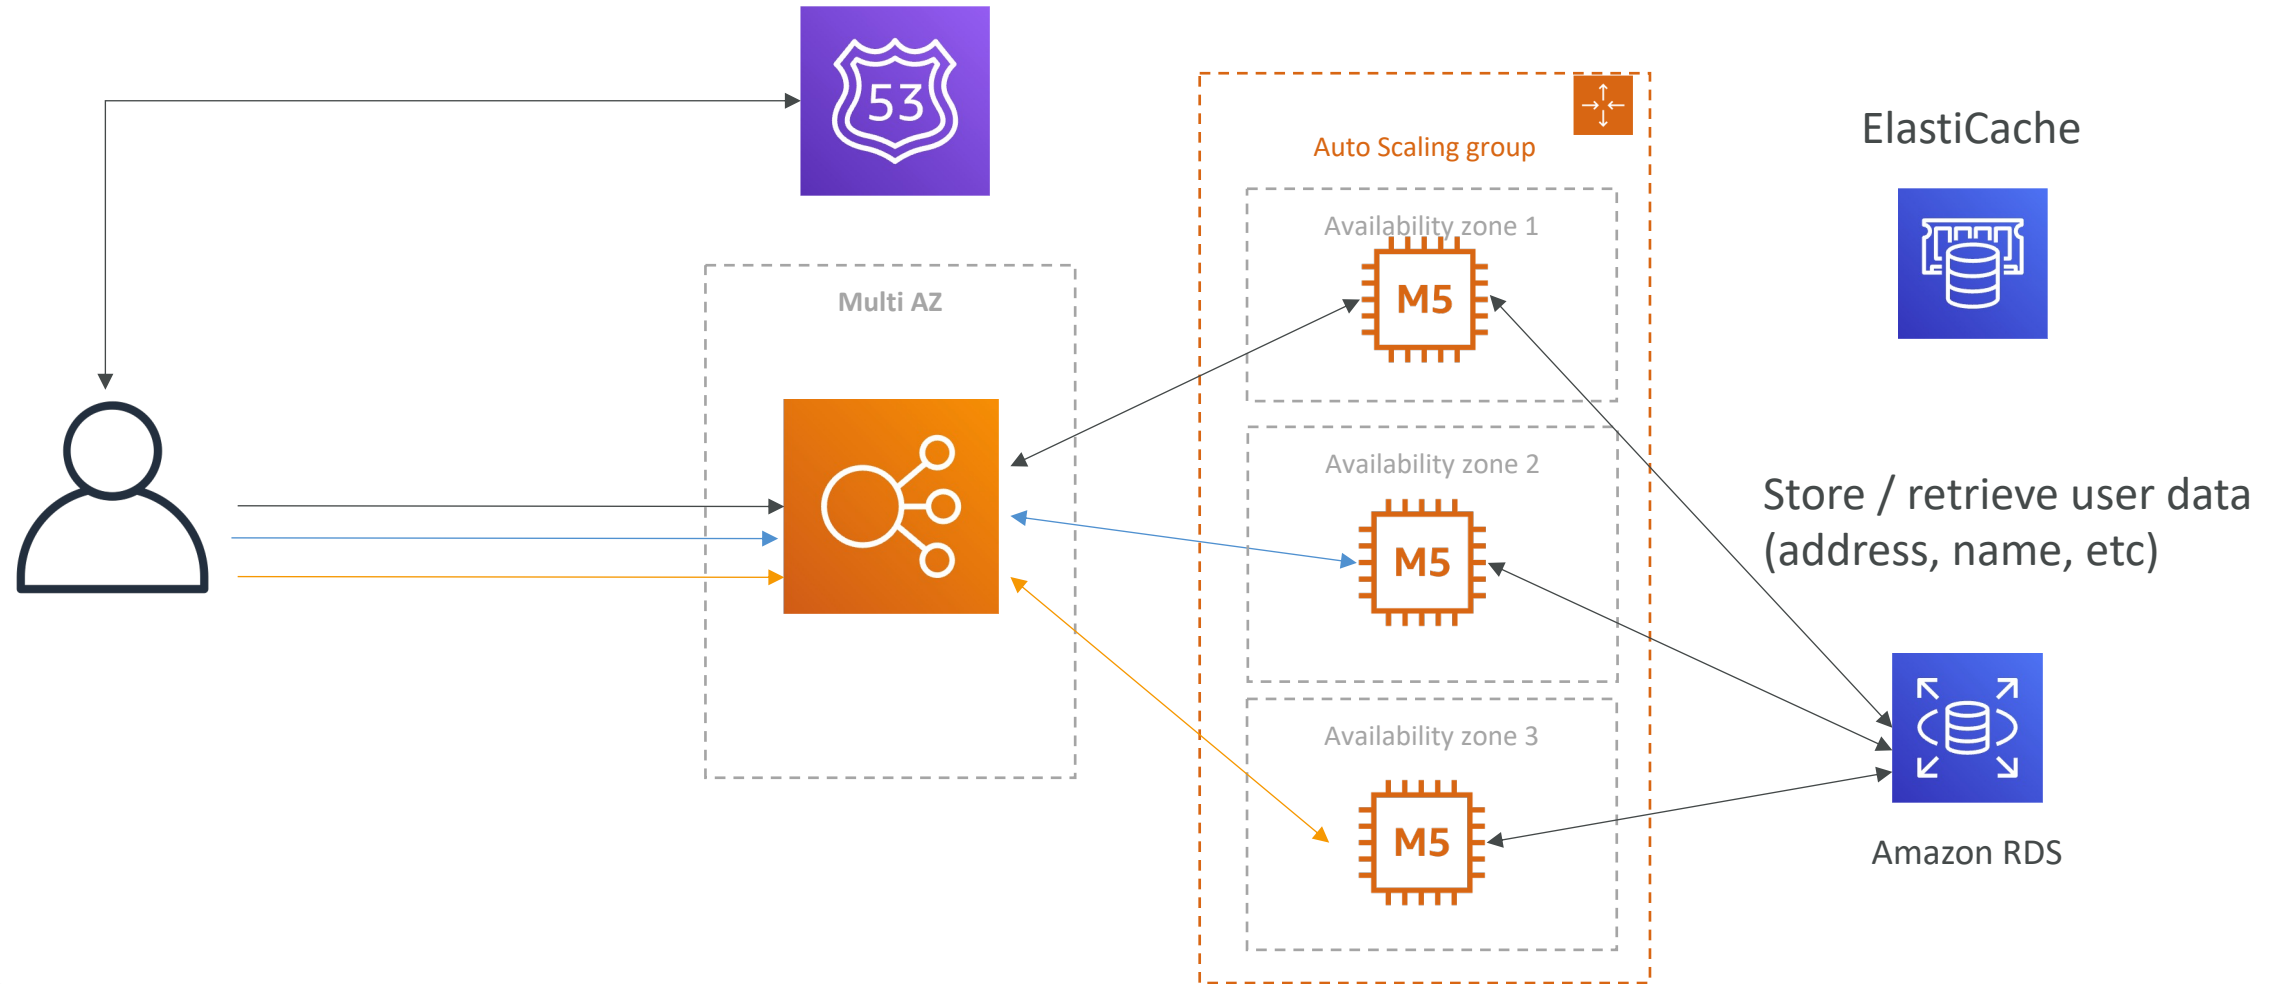
Stateful Web App: MyClothes.com

Introduce Server Session



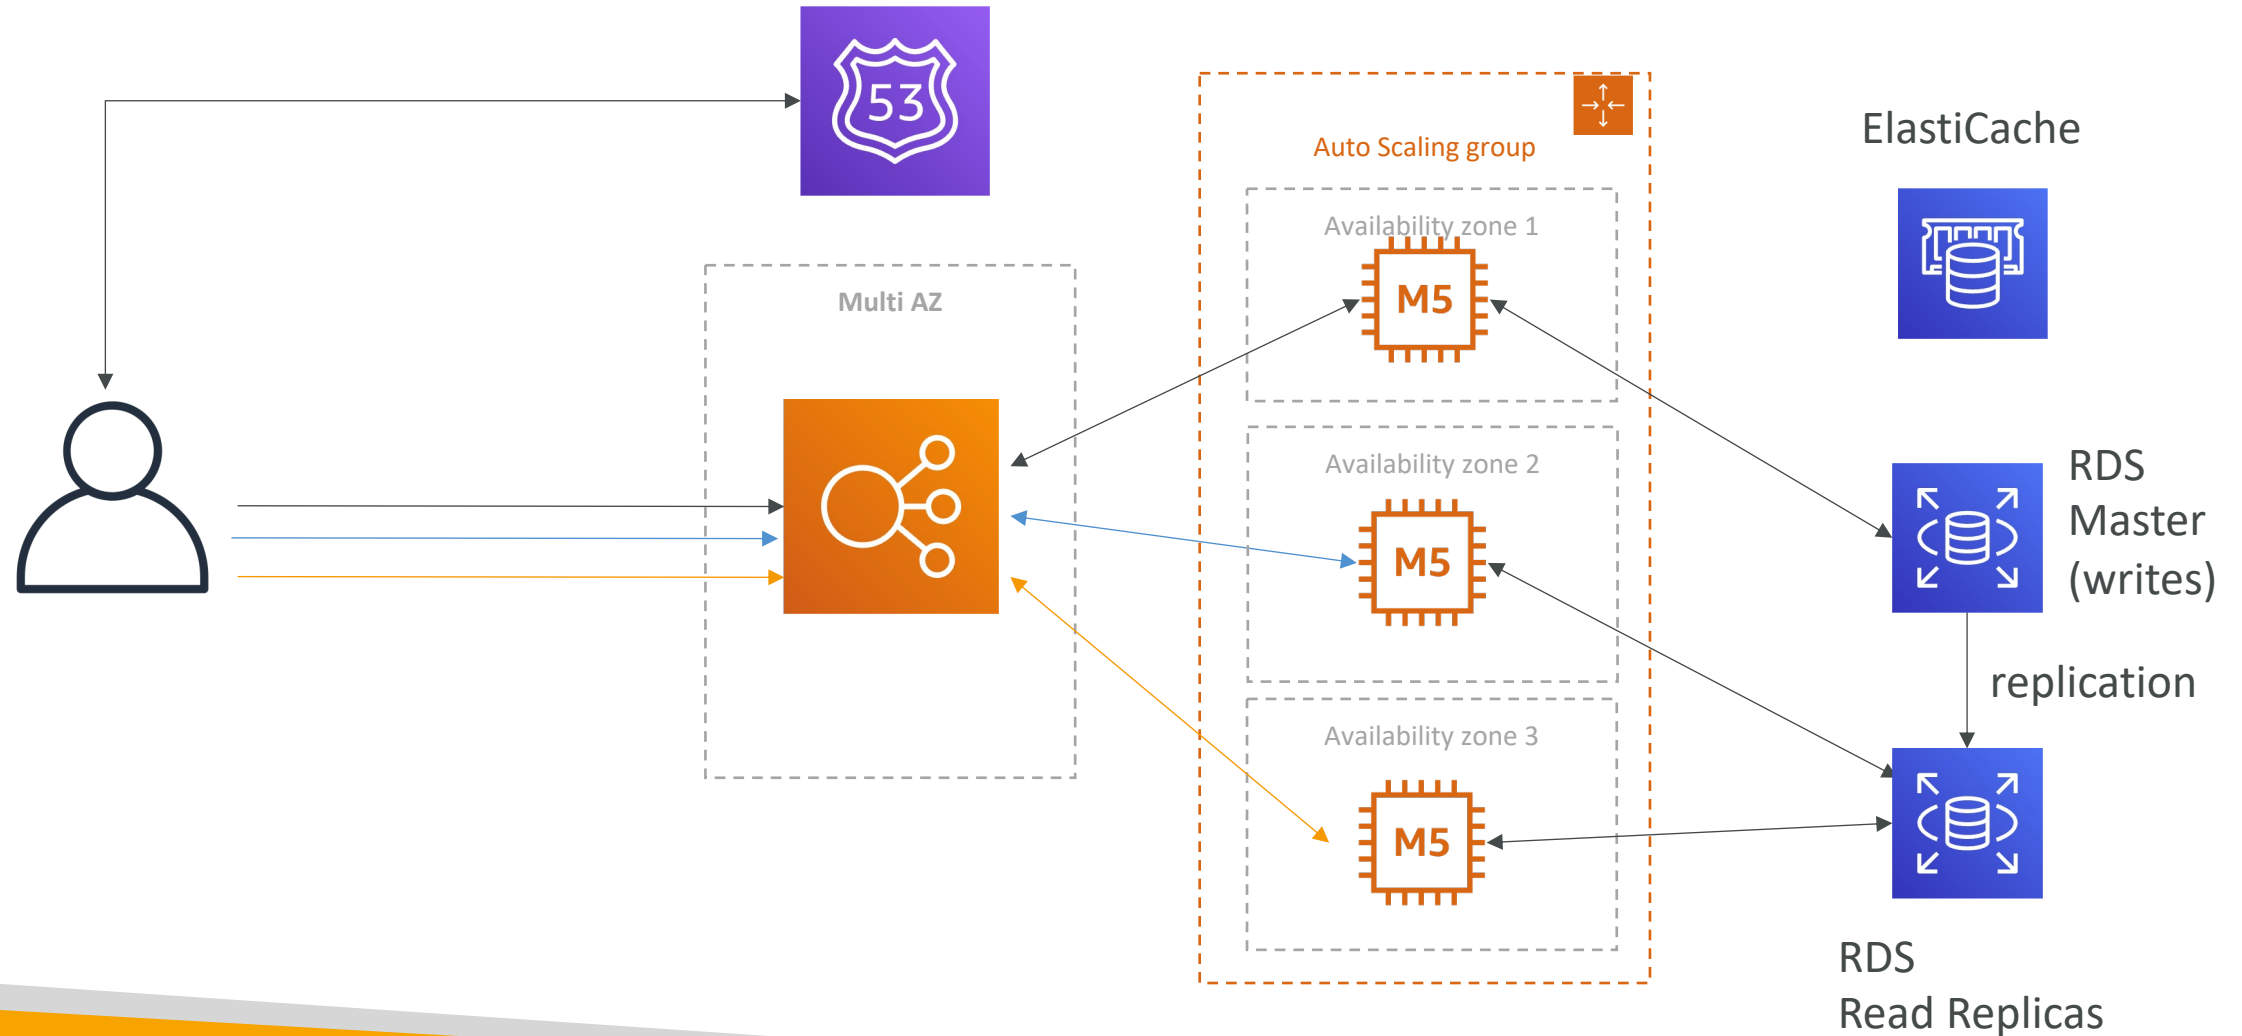
Stateful Web App: MyClothes.com

Storing User Data in a database



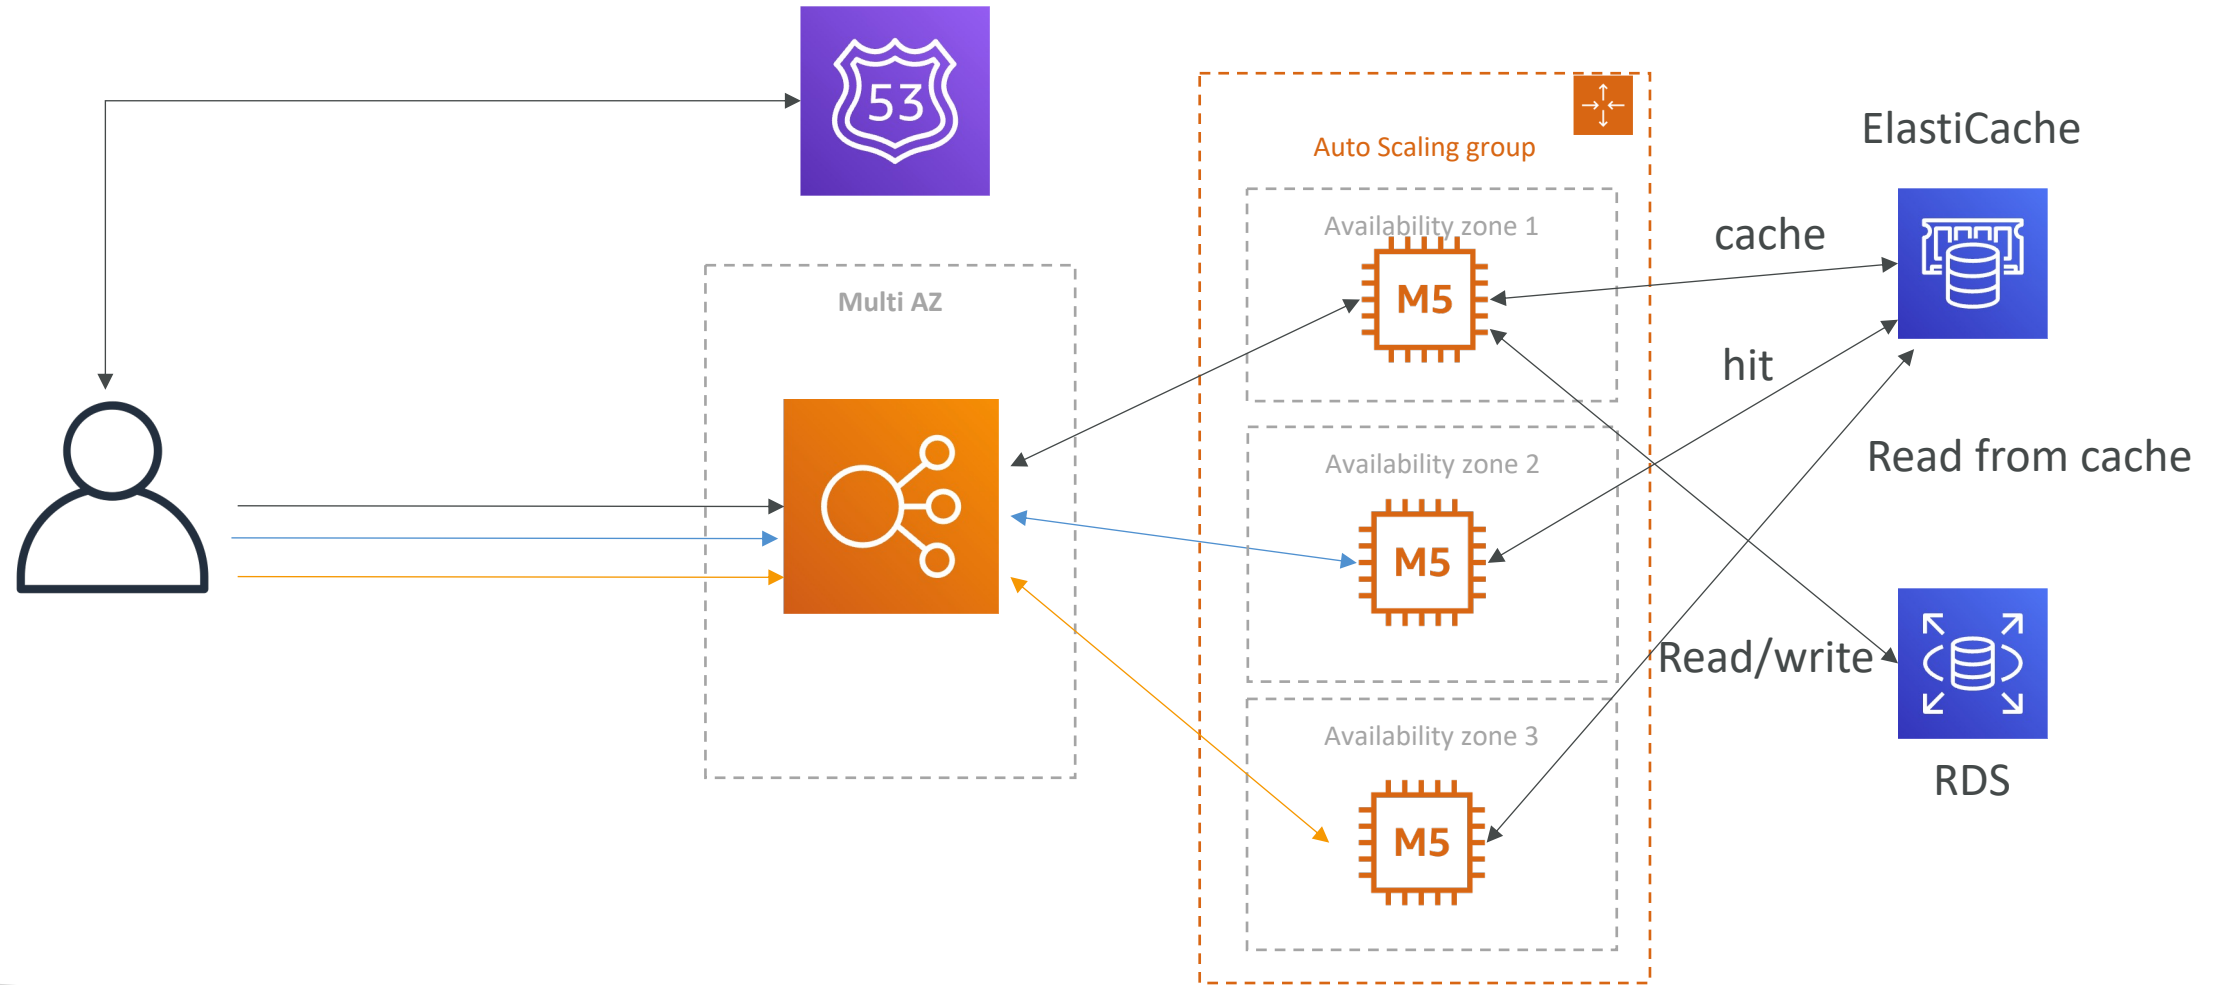
Stateful Web App: MyClothes.com

Scaling Reads



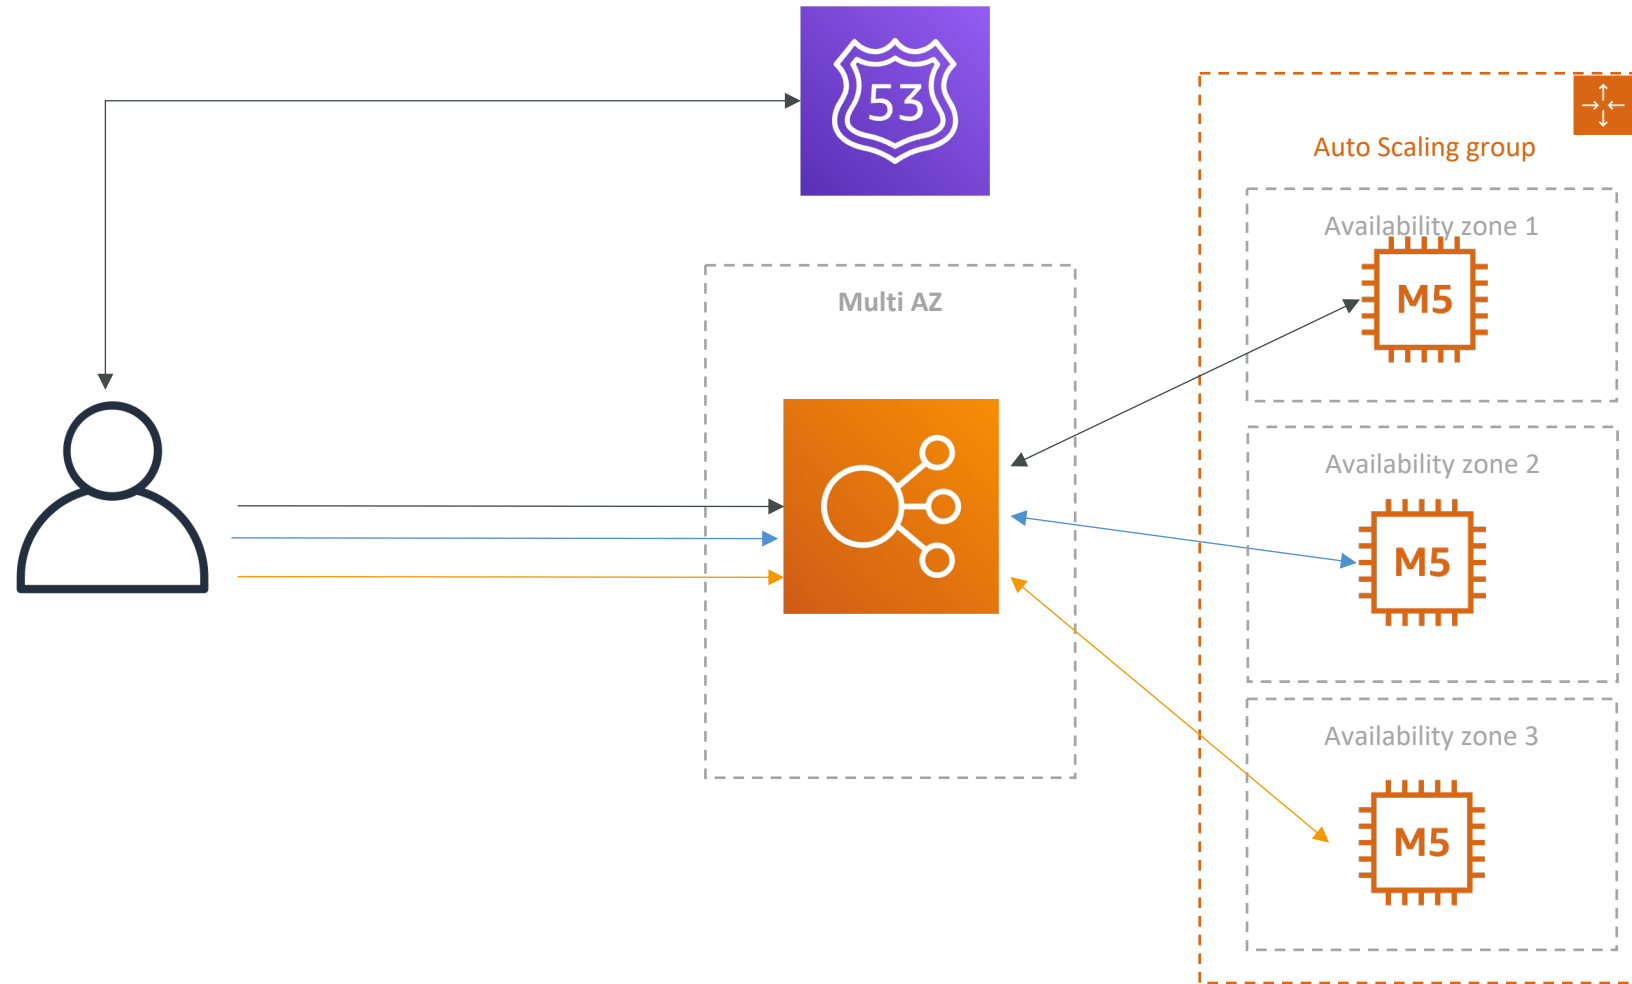
Stateful Web App: MyClothes.com

Scaling Reads (Alternative) – Lazy Loading

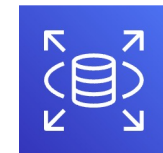
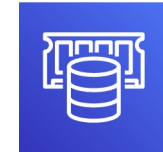


Stateful Web App: MyClothes.com

Multi AZ – Survive disasters



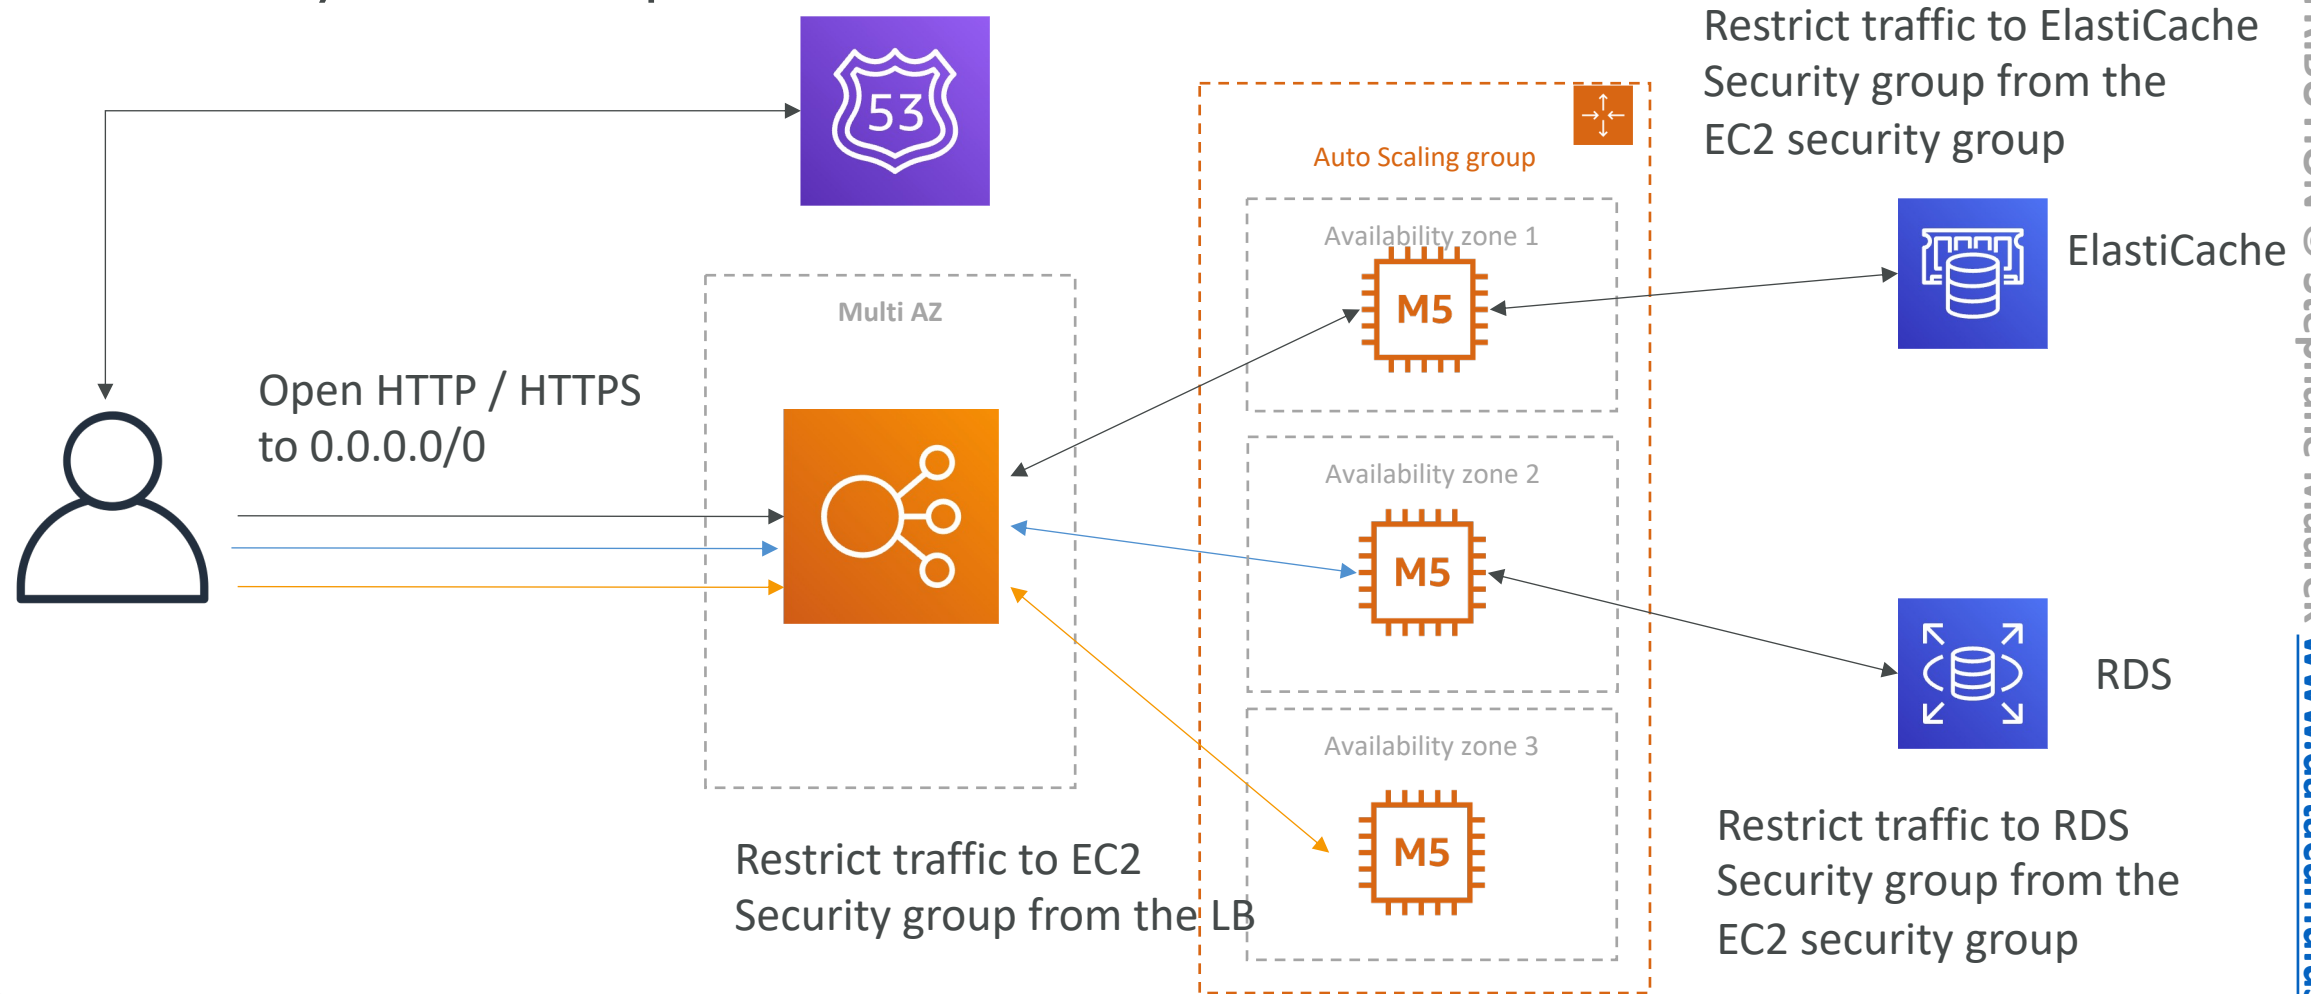
ElastiCache
Multi AZ



RDS
Multi AZ

Stateful Web App: MyClothes.com

Security Groups



In this lecture we've discussed...

3-tier architectures for web applications

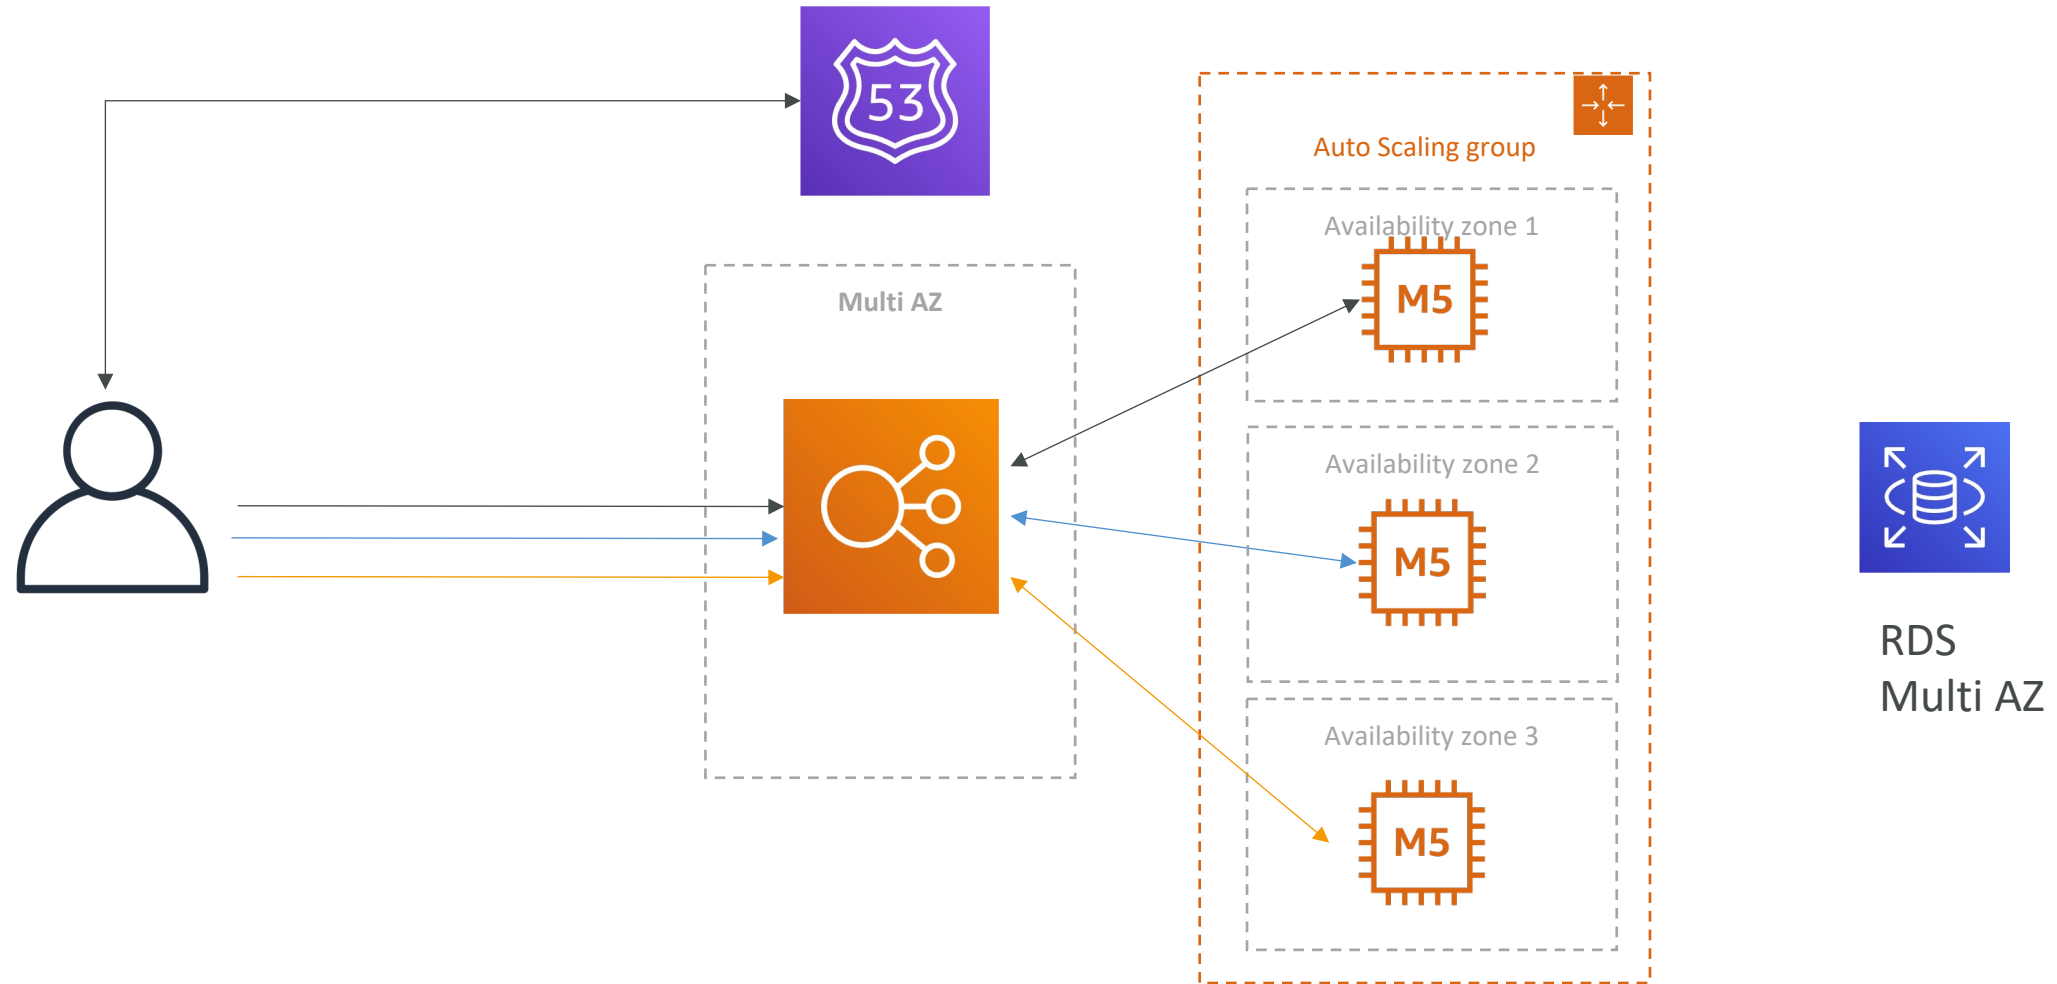
- ELB sticky sessions
- Web clients for storing cookies and making our web app stateless
- ElastiCache
 - For storing sessions (alternative: DynamoDB)
 - For caching data from RDS
 - Multi AZ
- RDS
 - For storing user data
 - Read replicas for scaling reads
 - Multi AZ for disaster recovery
- Tight Security with security groups referencing each other

Stateful Web App: MyWordPress.com

- We are trying to create a fully scalable WordPress website
- We want that website to access and correctly display picture uploads
- Our user data, and the blog content should be stored in a MySQL database.
- Let's see how we can achieve this!

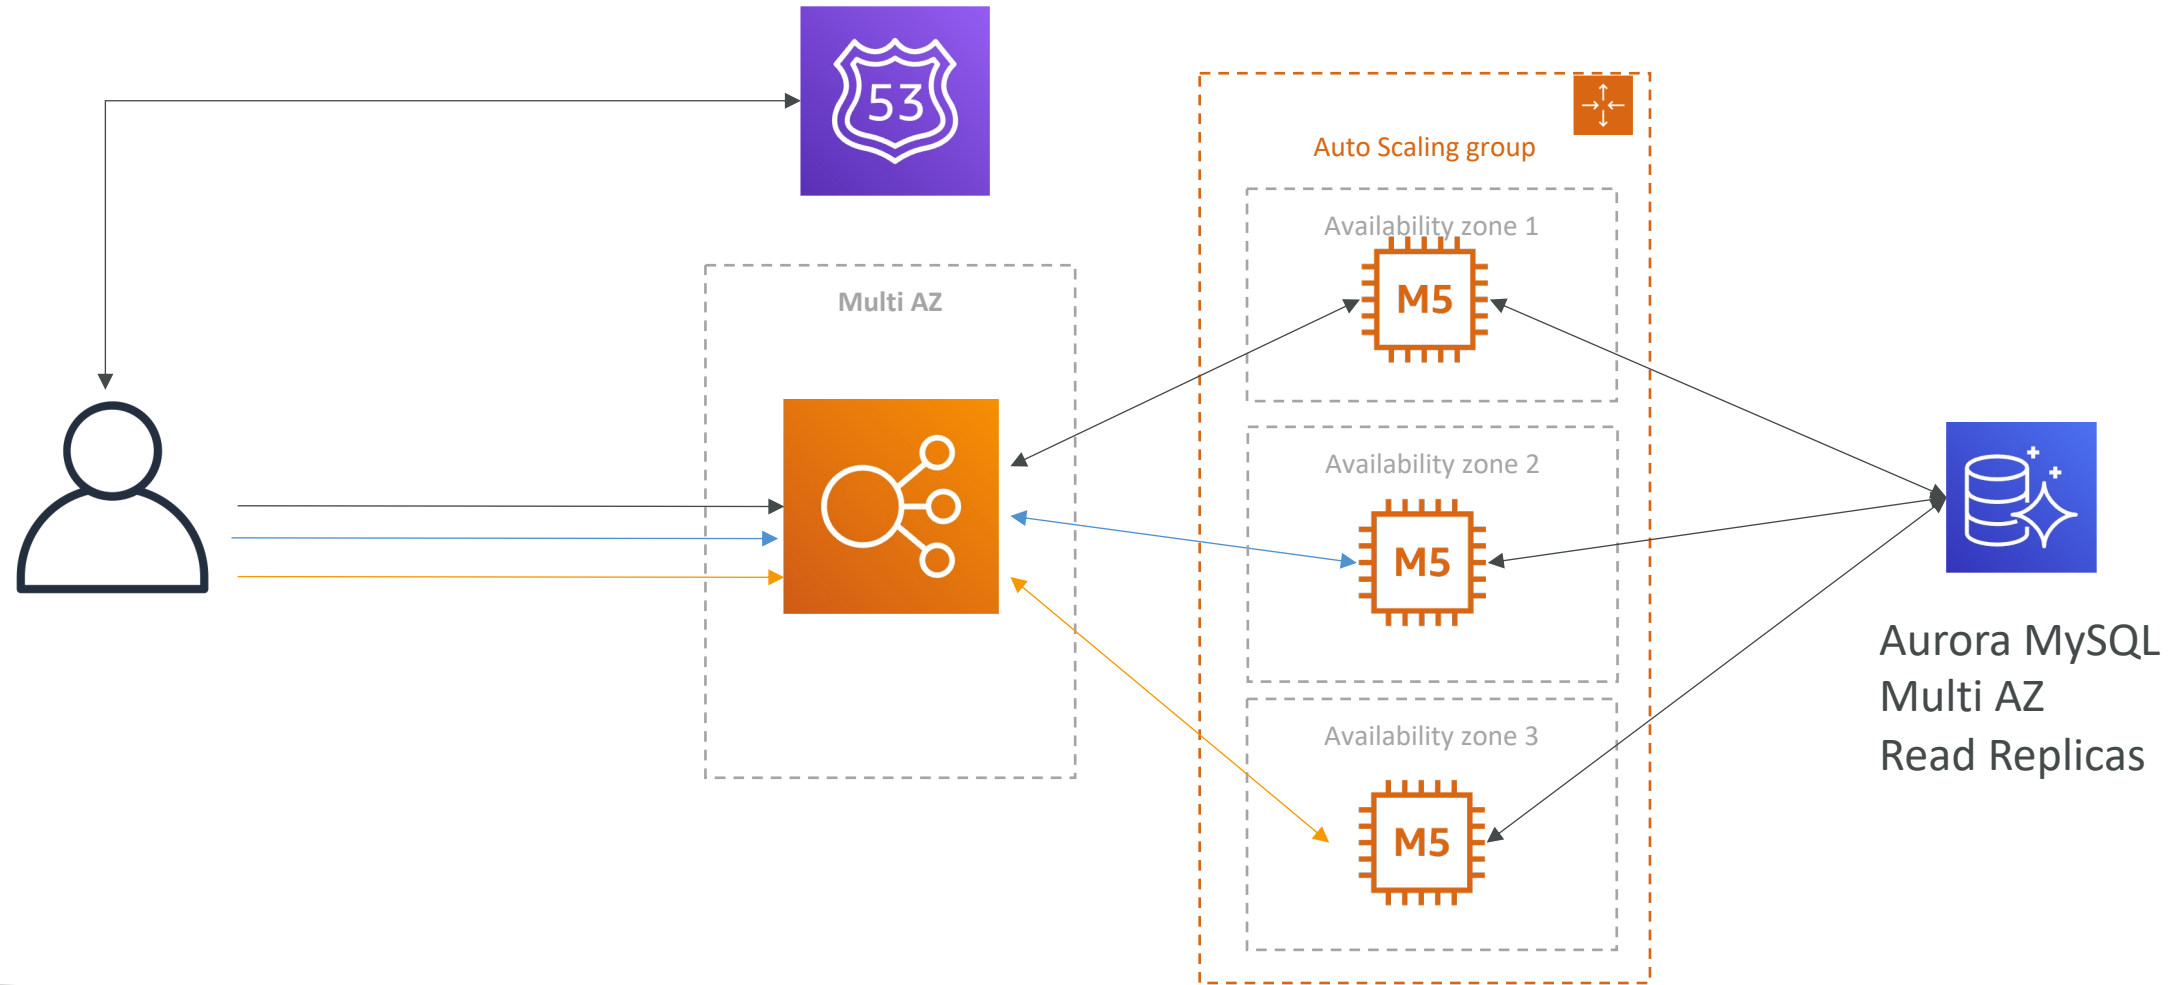
Stateful Web App: MyWordPress.com

RDS layer



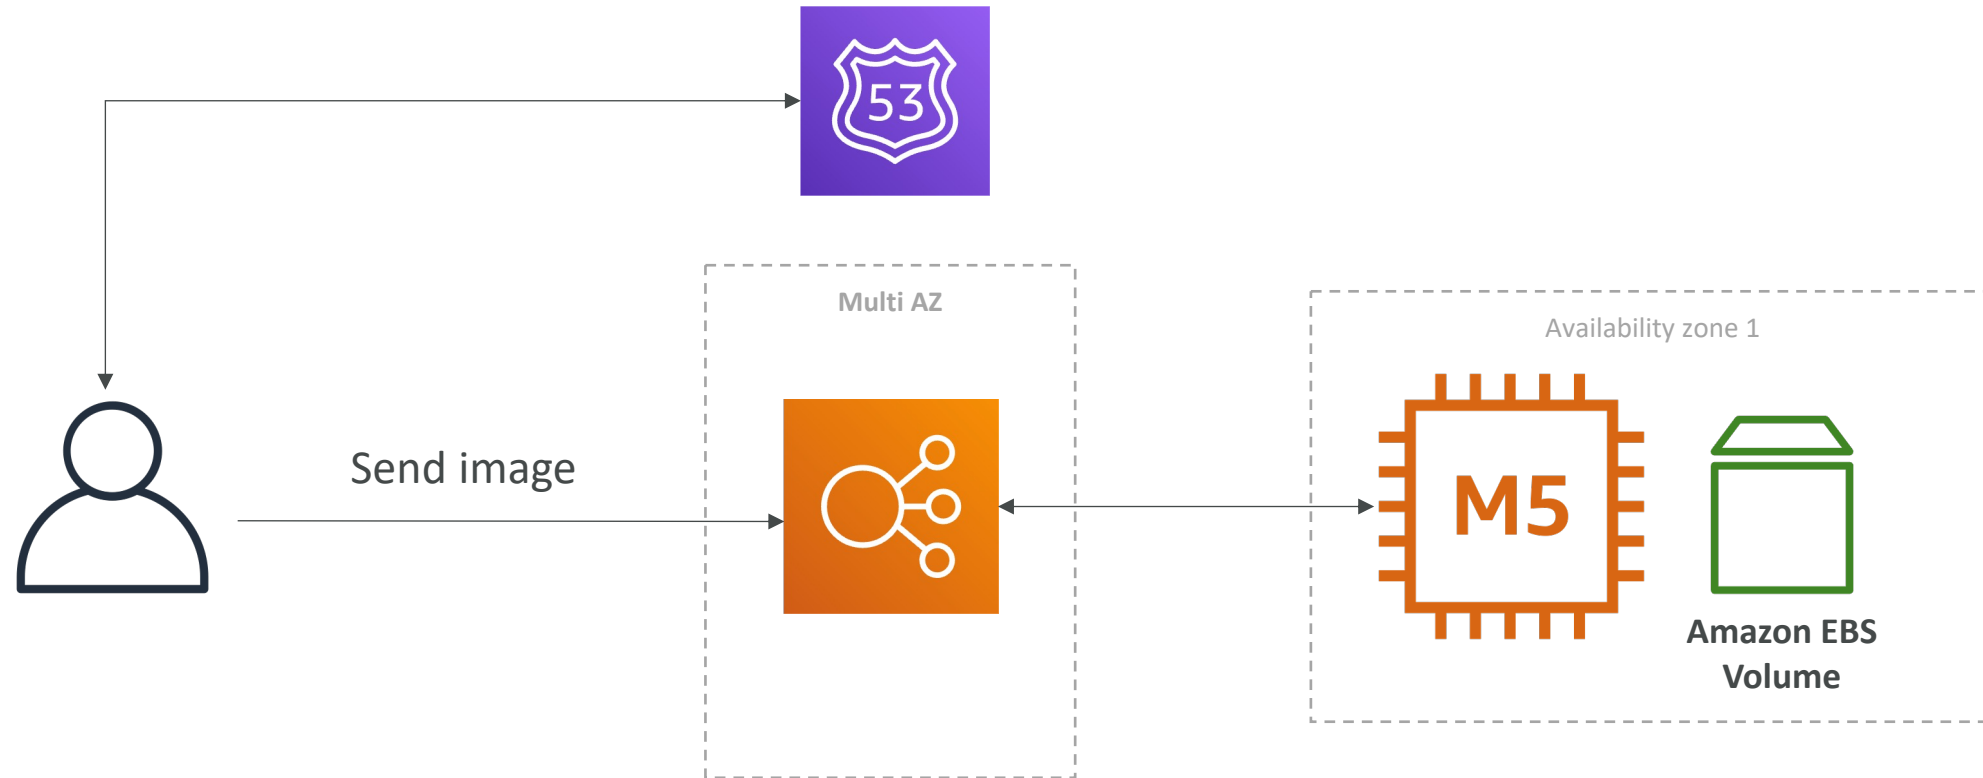
Stateful Web App: MyWordPress.com

Scaling with Aurora: Multi AZ & Read Replicas



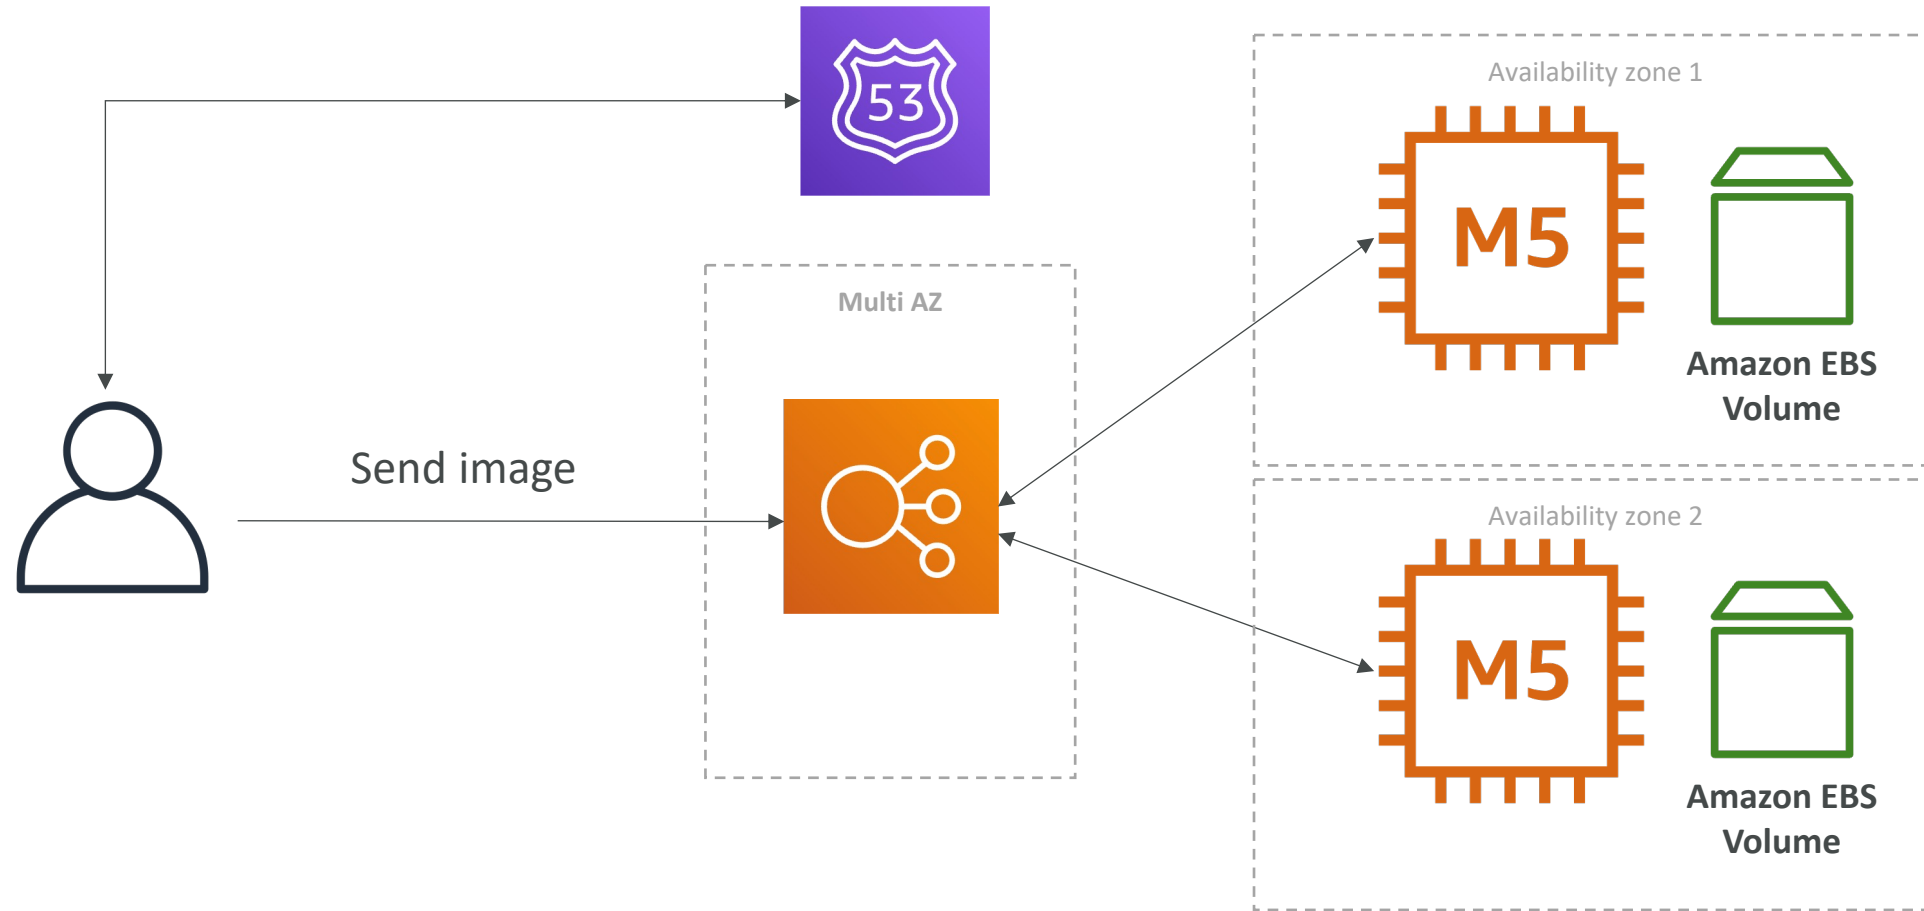
Stateful Web App: MyWordPress.com

Storing images with EBS



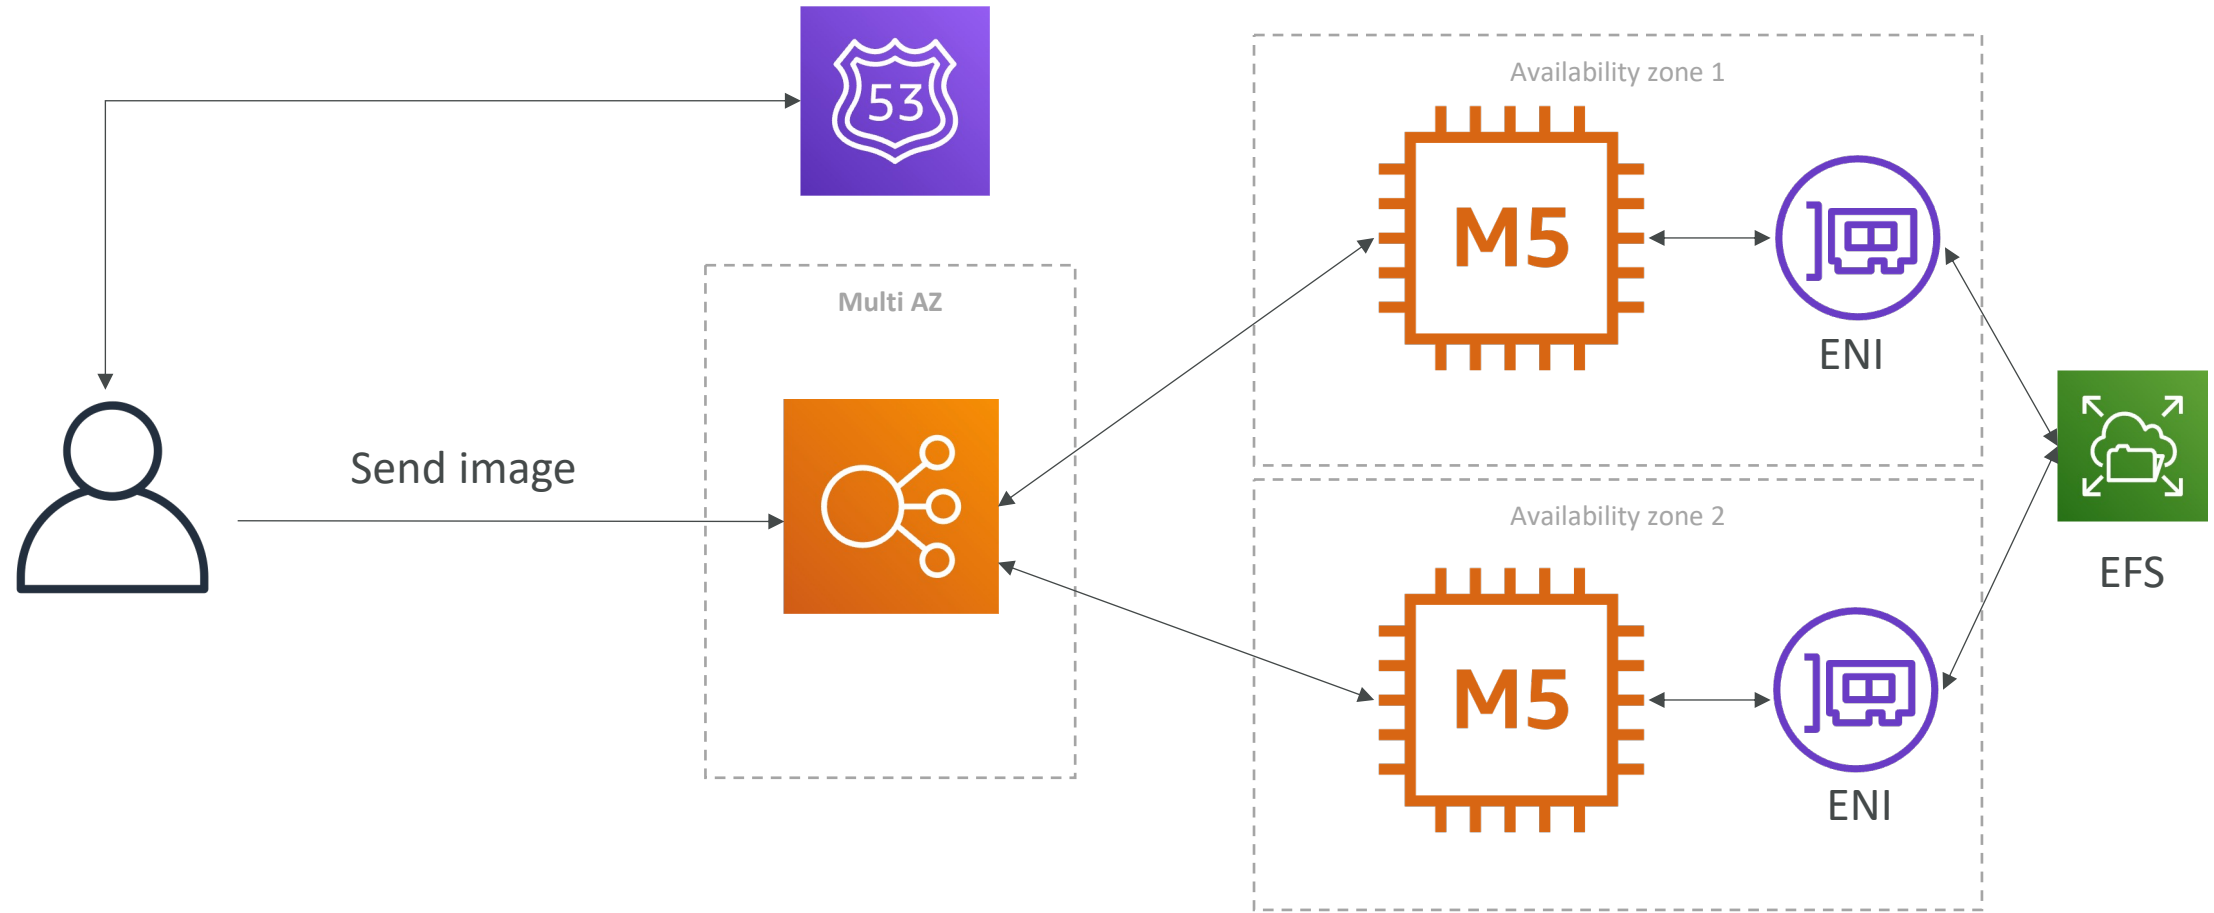
Stateful Web App: MyWordPress.com

Storing images with EBS



Stateful Web App: MyWordPress.com

Storing images with EFS



In this lecture we've discussed...

- Aurora Database to have easy Multi-AZ and Read-Replicas
- Storing data in EBS (single instance application)
- Vs Storing data in EFS (distributed application)

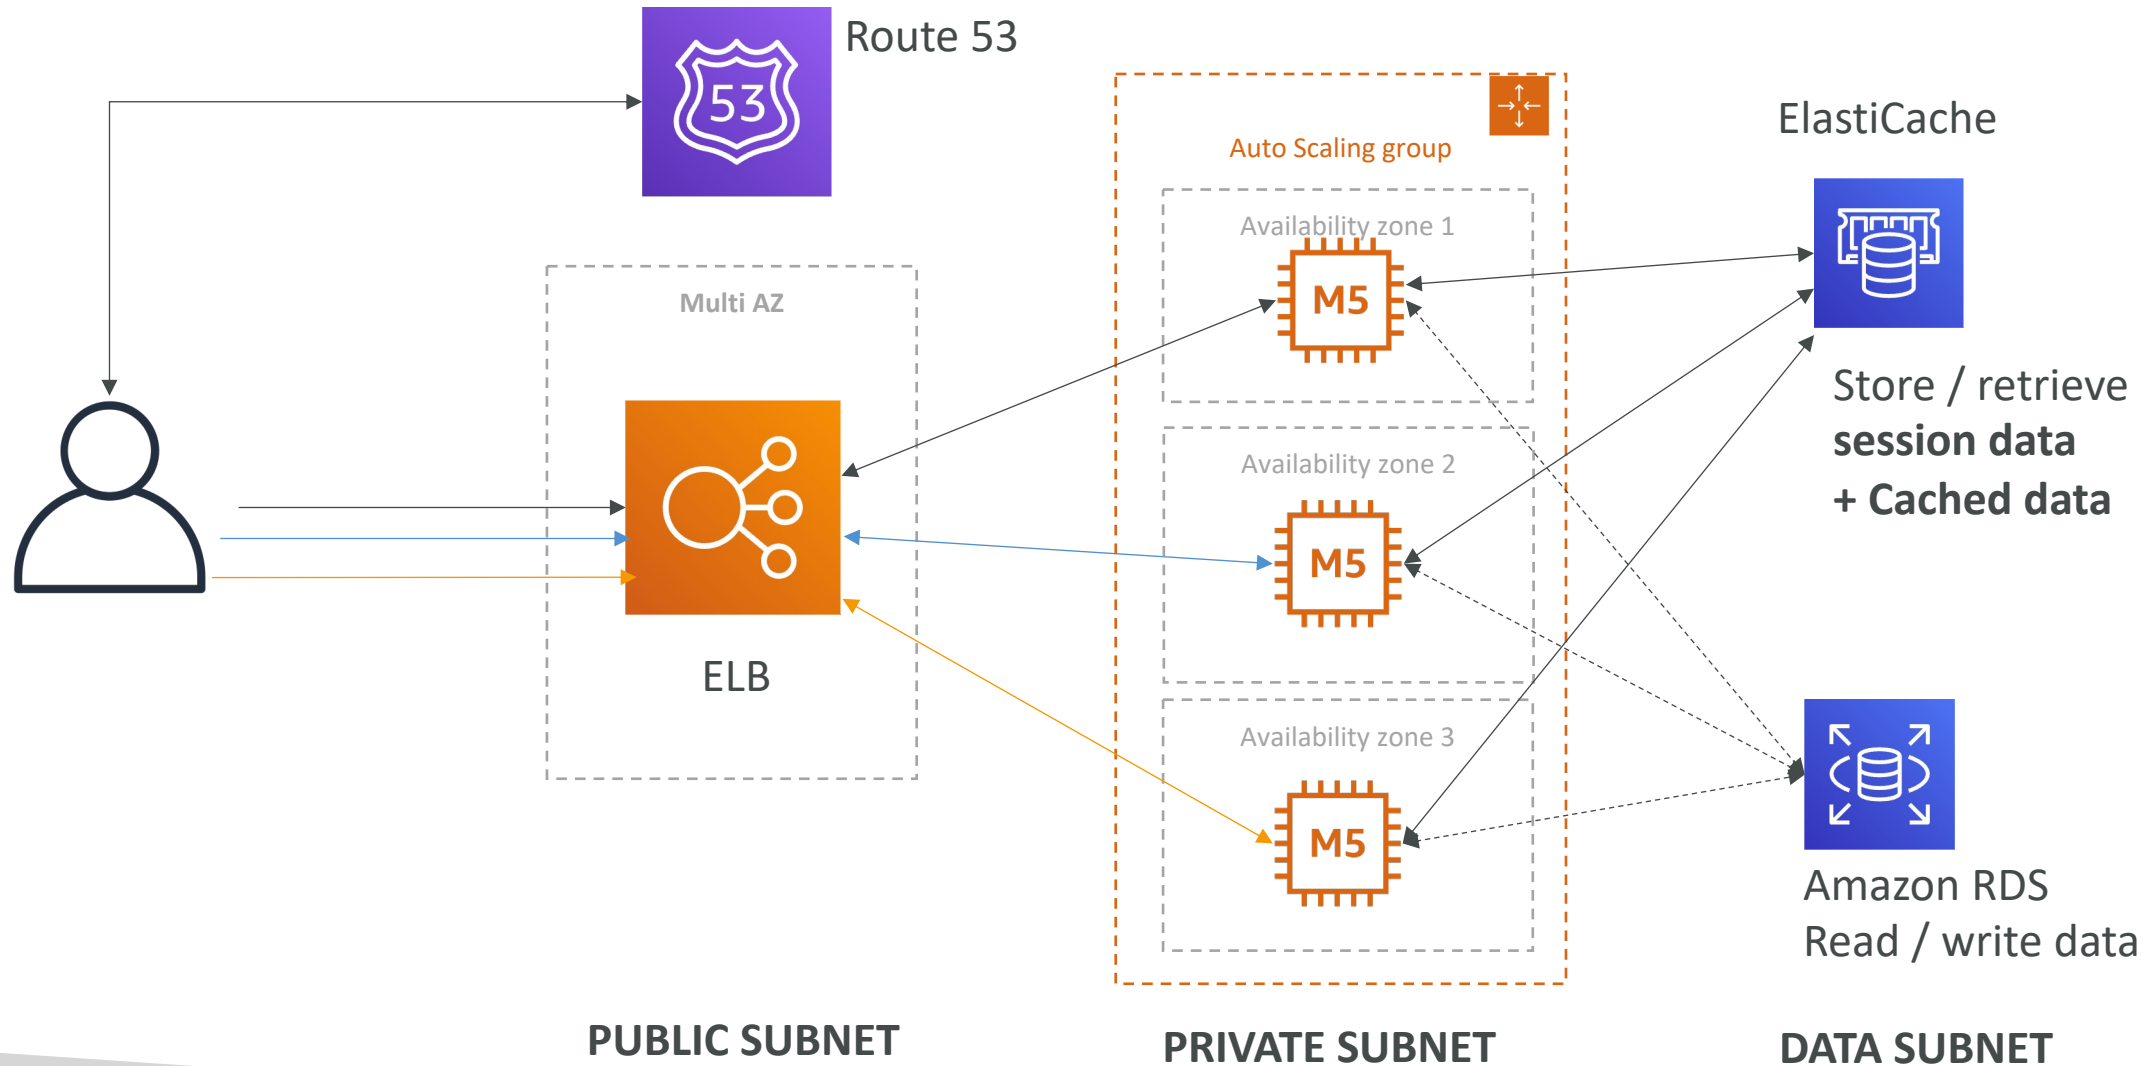
Instantiating Applications quickly

- When launching a full stack (EC2, EBS, RDS), it can take time to:
 - Install applications
 - Insert initial (or recovery) data
 - Configure everything
 - Launch the application
- We can take advantage of the cloud to speed that up!

Instantiating Applications quickly

- EC2 Instances:
 - **Use a Golden AMI:** Install your applications, OS dependencies etc.. beforehand and launch your EC2 instance from the Golden AMI
 - **Bootstrap using User Data:** For dynamic configuration, use User Data scripts
 - **Hybrid:** mix Golden AMI and User Data (Elastic Beanstalk)
- RDS Databases:
 - Restore from a snapshot: the database will have schemas and data ready!
- EBS Volumes:
 - Restore from a snapshot: the disk will already be formatted and have data!

Typical architecture: Web App 3-tier



Developer problems on AWS

- Managing infrastructure
 - Deploying Code
 - Configuring all the databases, load balancers, etc
 - Scaling concerns
-
- Most web apps have the same architecture (ALB + ASG)
 - All the developers want is for their code to run!
 - Possibly, consistently across different applications and environments

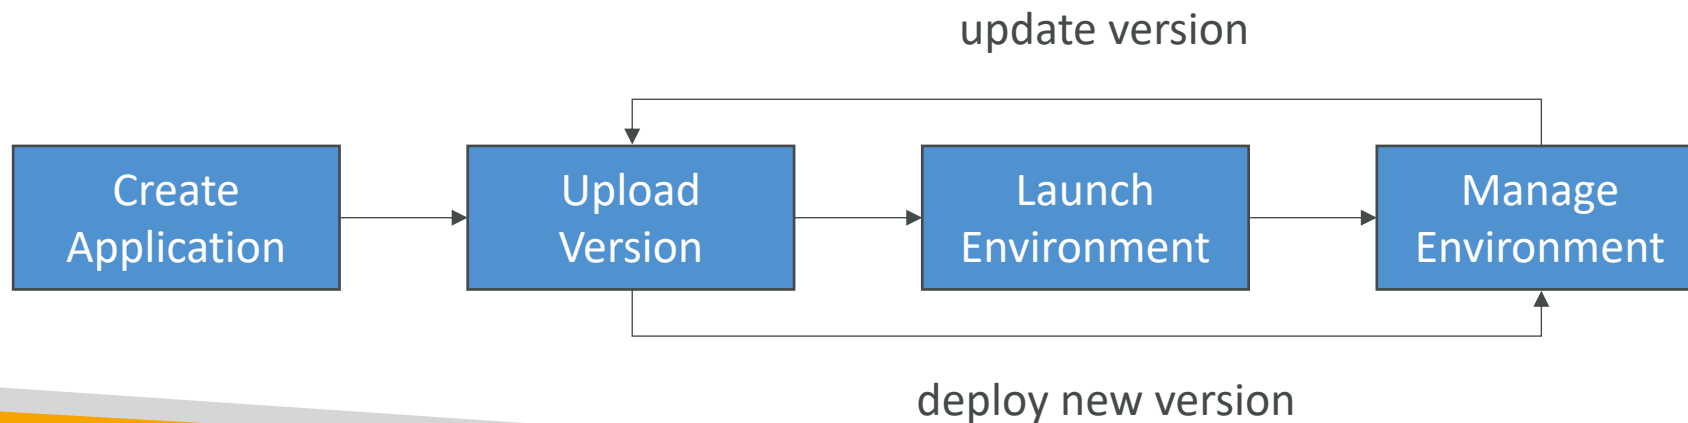


Elastic Beanstalk – Overview

- Elastic Beanstalk is a developer centric view of deploying an application on AWS
- It uses all the components we've seen before: EC2, ASG, ELB, RDS, ...
- Managed service
 - Automatically handles capacity provisioning, load balancing, scaling, application health monitoring, instance configuration, ...
 - Just the application code is the responsibility of the developer
- We still have full control over the configuration
- Beanstalk is free but you pay for the underlying instances

Elastic Beanstalk – Components

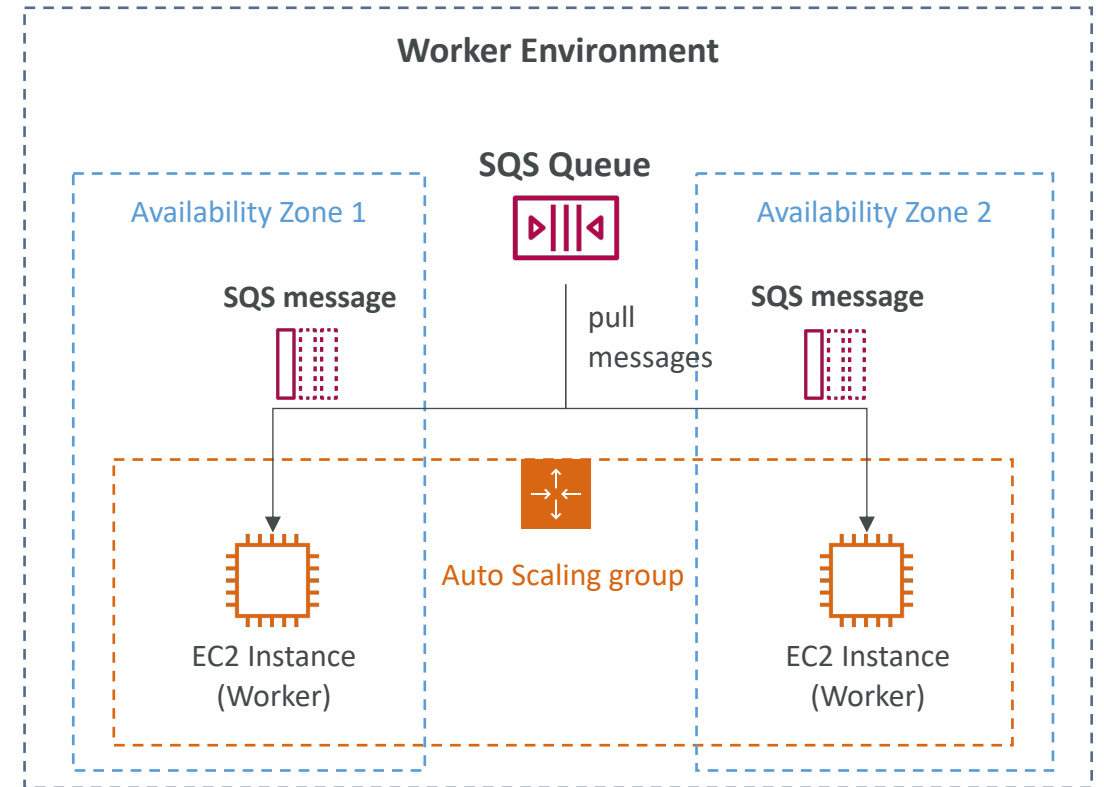

- **Application:** collection of Elastic Beanstalk components (environments, versions, configurations, ...)
- **Application Version:** an iteration of your application code
- **Environment**
 - Collection of AWS resources running an application version (only one application version at a time)
 - **Tiers:** Web Server Environment Tier & Worker Environment Tier
 - You can create multiple environments (dev, test, prod, ...)



Elastic Beanstalk – Supported Platforms

- Go
- Java SE
- Java with Tomcat
- .NET Core on Linux
- .NET on Windows Server
- Node.js
- PHP
- Python
- Ruby
- Packer Builder
- Single Container Docker
- Multi-container Docker
- Preconfigured Docker

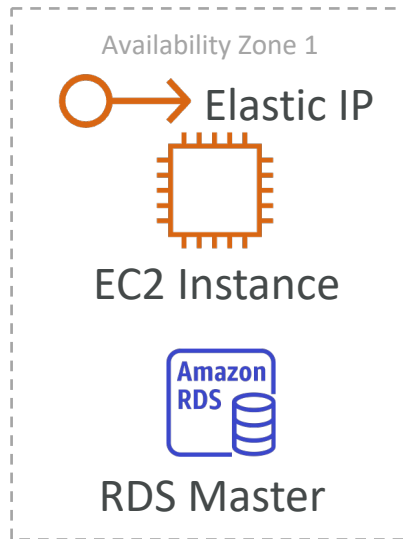
Web Server Tier vs. Worker Tier



- Scale based on the number of SQS messages
- Can push messages to SQS queue from another Web Server Tier

Elastic Beanstalk Deployment Modes

Single Instance
Great for dev



High Availability with Load Balancer
Great for prod

